

# Program Analysis and Finding Vulnerabilities

Vitaly Shmatikov

# Language-Based Approaches

---

(More) **type-safe** languages prevent some vulns by design

- “A language is type-safe if the only operations that can be performed on data in the language are those sanctioned by the type of the data.”
- Traditionally less performance

New generation of safer high-performance languages:

- Rust (Mozilla), Swift (Apple), Go (Google)

Efforts to improve security of unsafe languages

Safe pointer libraries in C / C++

Coding standards, defensive programming, unit testing

# Software Engineering Approaches

---

Organize software lifecycle around security

Require use of organizational and software tools to improve security outcomes

Microsoft security development lifecycle (SDL):

Training	Manage risk of third-party components
Design security requirements	Use approved tools
Metrics & compliance reporting	Static analysis security testing
Threat modeling	Dynamic analysis security testing
Establish design requirements	Penetration testing
Define & use crypto standards	Incident response

# Most Software Very, Very Complex

---

In a Nutshell, Apache HTTP Server...

... has had 39,732 commits made by 125 contributors

representing 1,494,342 lines of code

... is mostly written in C  
with an average number of source code comments

Linux kernel v.4.1:

~19.5 million lines of code


14,000 developers contributing

OpenSSL:

~608,000 lines of code

572 developers contributing

# Remember Heartbleed?



OpenSSL implements TLS, used in Apache and Nginx  
March 2014: researchers discover vulnerability in  
the OpenSSL implementation of TLS heartbeat

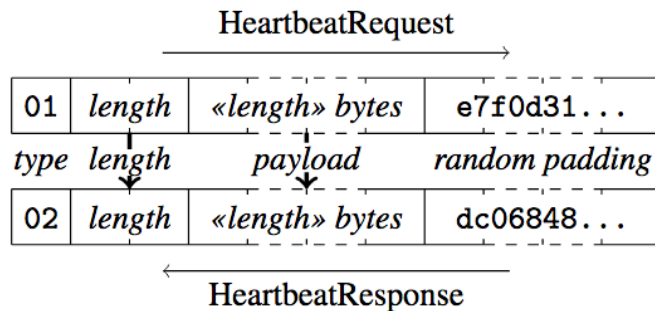


Figure 1: **Heartbeat Protocol.** Heartbeat requests include user data and random padding. The receiving peer responds by echoing back the data in the initial request along with its own padding.

[Durumeric et al. 2014]

# TLS Heartbeat

A way to keep TLS connection alive  
without constantly transferring data

If you are alive, send me  
this 5-letter word: “xyzzy”

“xyzzy”

C

Per RFC 6520:

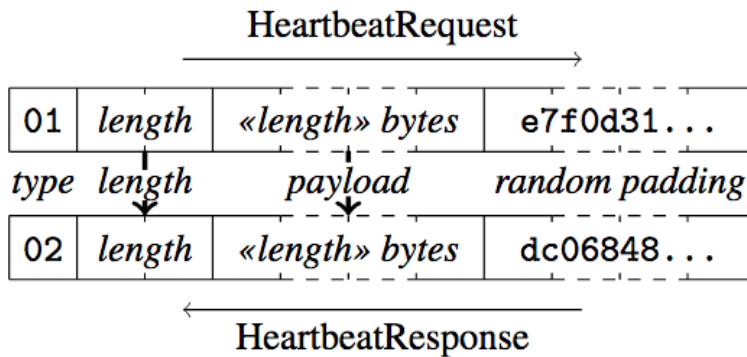
```
struct {  
    HeartbeatMessageType type;  
    uint16 payload_length;  
    opaque payload[HeartbeatMessage.payload_length];  
    opaque padding[padding_length];  
} HeartbeatMessage;
```

OpenSSL omitted to  
check that this value  
matches the actual length  
of the heartbeat message

S



# Heartbleed



Buffer overread vulnerability  
Copy up to almost  $2^{16}$  bytes  
of data from memory

```
1448 dtls1_process_heartbeat(SSL *s)
1449 {
1450     unsigned char *p = &s->s3->rrec.data[0], *pl;
1451     unsigned short hbtype;
1452     unsigned int payload;
1453     unsigned int padding = 16; /* Use minimum padding */
1454
1455     /* Read type and payload length first */
1456     hbtype = *p++;
1457     n2s(p, payload);
1458     pl = p;
1459
1460     if (s->msg_callback)
1461         s->msg_callback(0, s->version, TLS1_RT_HEARTBEAT,
1462             &s->s3->rrec.data[0], s->s3->rrec.length,
1463             s, s->msg_callback_arg);
1464
1465     if (hbtype == TLS1_HB_REQUEST)
1466     {
1467         unsigned char *buffer, *bp;
1468         int r;
1469
1470         /* Allocate memory for the response, size is 1 byte
1471          * message type, plus 2 bytes payload length, plus
1472          * payload, plus padding
1473          */
1474         buffer = OPENSSL_malloc(1 + 2 + payload + padding);
1475         bp = buffer;
1476
1477         /* Enter response type, length and copy payload */
1478         *bp++ = TLS1_HB_RESPONSE;
1479         s2n(payload, bp);
1480         memcpy(bp, pl, payload);
1481         bp += payload;
```

# Heartbleed Chronology

---

"I was doing laborious auditing of OpenSSL, going through the [Secure Sockets Layer] stack line by line"

Date	Event
03/21	Neel Mehta of Google discovers Heartbleed
03/21	Google patches OpenSSL on their servers
03/31	CloudFlare is privately notified and patches
04/01	Google notifies the OpenSSL core team
04/02	Codenomicon independently discovers Heartbleed
04/03	Codenomicon informs NCSC-FI
04/04	Akamai is privately notified and patches
04/05	Codenomicon purchases the <code>heartbleed.com</code> domain
04/06	OpenSSL notifies several Linux distributions
04/07	NCSC-FI notifies OpenSSL core team
04/07	OpenSSL releases version 1.0.1g and a security advisory
04/07	CloudFlare and Codenomicon disclose on Twitter
04/08	Al-Bassam scans the Alexa Top 10,000
04/09	University of Michigan begins scanning

[Durumeric et al. 2014]



# Scanning for Heartbleed

Internet scanning to determine vulnerability:

Send heartbeat request with zero length (indicates vulnerable system)

Web Server	Alexa Sites	Heartbeat Ext.	Vulnerable
Apache	451,270 (47.3%)	95,217 (58.4%)	28,548 (64.4%)
Nginx	182,379 (19.1%)	46,450 (28.5%)	11,185 (25.2%)
Microsoft IIS	96,259 (10.1%)	637 (0.4%)	195 (0.4%)
Litespeed	17,597 (1.8%)	6,838 (4.2%)	1,601 (3.6%)
Other	76,817 (8.1%)	5,383 (3.3%)	962 (2.2%)
Unknown	129,006 (13.5%)	8,545 (5.2%)	1,833 (4.1%)

[Durumeric et al. 2014]

# Scanning for Heartbleed

Internet scanning to determine vulnerability:

Send heartbeat request with zero length (indicates vulnerable system)

Site	Vuln.	Site	Vuln.	Site	Vuln.
Google	Yes	Bing	No	Wordpress	Yes
Facebook	No	Pinterest	Yes	Huff. Post	?
Youtube	Yes	Blogspot	Yes	ESPN	?
Yahoo	Yes	Go.com	?	Reddit	Yes
Amazon	No	Live	No	Netflix	Yes
Wikipedia	Yes	CNN	?	MSN.com	No
LinkedIn	No	Instagram	Yes	Weather.com	?
eBay	No	Paypal	No	IMDB	No
Twitter	No	Tumblr	Yes	Apple	No
Craigslist	?	Imgur	Yes	Yelp	?


[Durumeric et al. 2014]

# Disassembly and Decompiling


---

Heartbleed discovered by direct C code inspection  
What if you only have the binary?

Normal compilation  
process


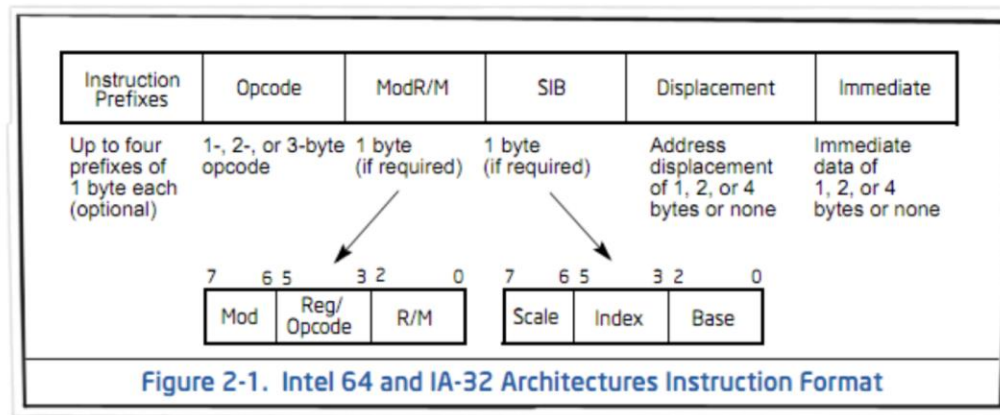


What if we start with  
binary?



Disassembler  
(gdb, IDA Pro, OllyDebug)

# Disassembly



## What type of vulnerability might this be?

```
Dump of assembler code for function main:
0x08048434 <main+0>:  push    %ebp
0x08048435 <main+1>:  mov     %esp,%ebp
0x08048437 <main+3>:  and     $0xffffffff0,%esp
0x0804843a <main+6>:  sub     $0x20,%esp
0x0804843d <main+9>:  movl    $0xf8, (%esp)
0x08048444 <main+16>: call     0x8048364 <malloc@plt>
0x08048449 <main+21>: mov     %eax,0x14(%esp)
0x0804844d <main+25>: movl    $0xf8, (%esp)
0x08048454 <main+32>: call     0x8048364 <malloc@plt>
0x08048459 <main+37>: mov     %eax,0x18(%esp)
0x0804845d <main+41>: mov     0x14(%esp),%eax
0x08048461 <main+45>: mov     %eax, (%esp)
0x08048464 <main+48>: call     0x8048354 <free@plt>
0x08048469 <main+53>: mov     0x18(%esp),%eax
0x0804846d <main+57>: mov     %eax, (%esp)
0x08048470 <main+60>: call     0x8048354 <free@plt>
0x08048475 <main+65>: movl    $0x200, (%esp)
0x0804847c <main+72>: call     0x8048364 <malloc@plt>
0x08048481 <main+77>: mov     %eax,0x1c(%esp)
0x08048485 <main+81>:
0x08048488 <main+84>:
0x0804848b <main+87>:
0x0804848d <main+89>:
0x08048495 <main+97>:
0x08048499 <main+101>:
0x0804849d <main+105>:
0x080484a0 <main+108>:
0x080484a5 <main+113>:
0x080484a9 <main+117>:
0x080484ac <main+120>: call     0x8048354 <free@plt>
0x080484b1 <main+125>: mov     0x1c(%esp),%eax
0x080484b5 <main+129>: mov     %eax, (%esp)
0x080484b8 <main+132>: call     0x8048354 <free@plt>
0x080484bd <main+137>: leave
0x080484be <main+138>: ret
End of assembler dump.
(gdb)
```

Double-free vulnerability

Exploit can trick heap management software into writing adversary-controlled value to adversary-controlled address

```
main( int argc, char* argv[] ) {
    char* b1;
    char* b2;
    char* b3;
```


```
    if( argc != 3 ) then return 0;
    if( atoi(argv[2]) != 31337 )
        complicatedFunction();
```

```
    else {
        b1 = (char*)malloc(248);
        b2 = (char*)malloc(248);
        free(b1);
        free(b2);
        b3 = (char*)malloc(512);
        strncpy( b3, argv[1], 511 );
        free(b2);
        free(b3);
```

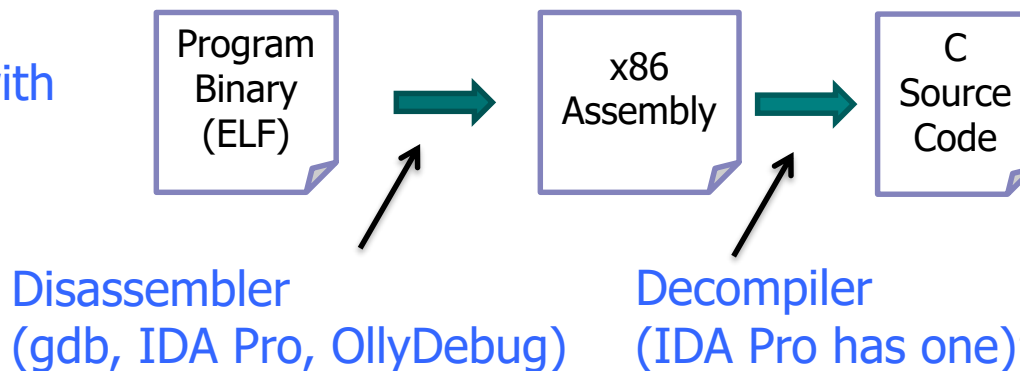
```
    }
}
```

# Disassembly and Decompiling

Normal compilation process

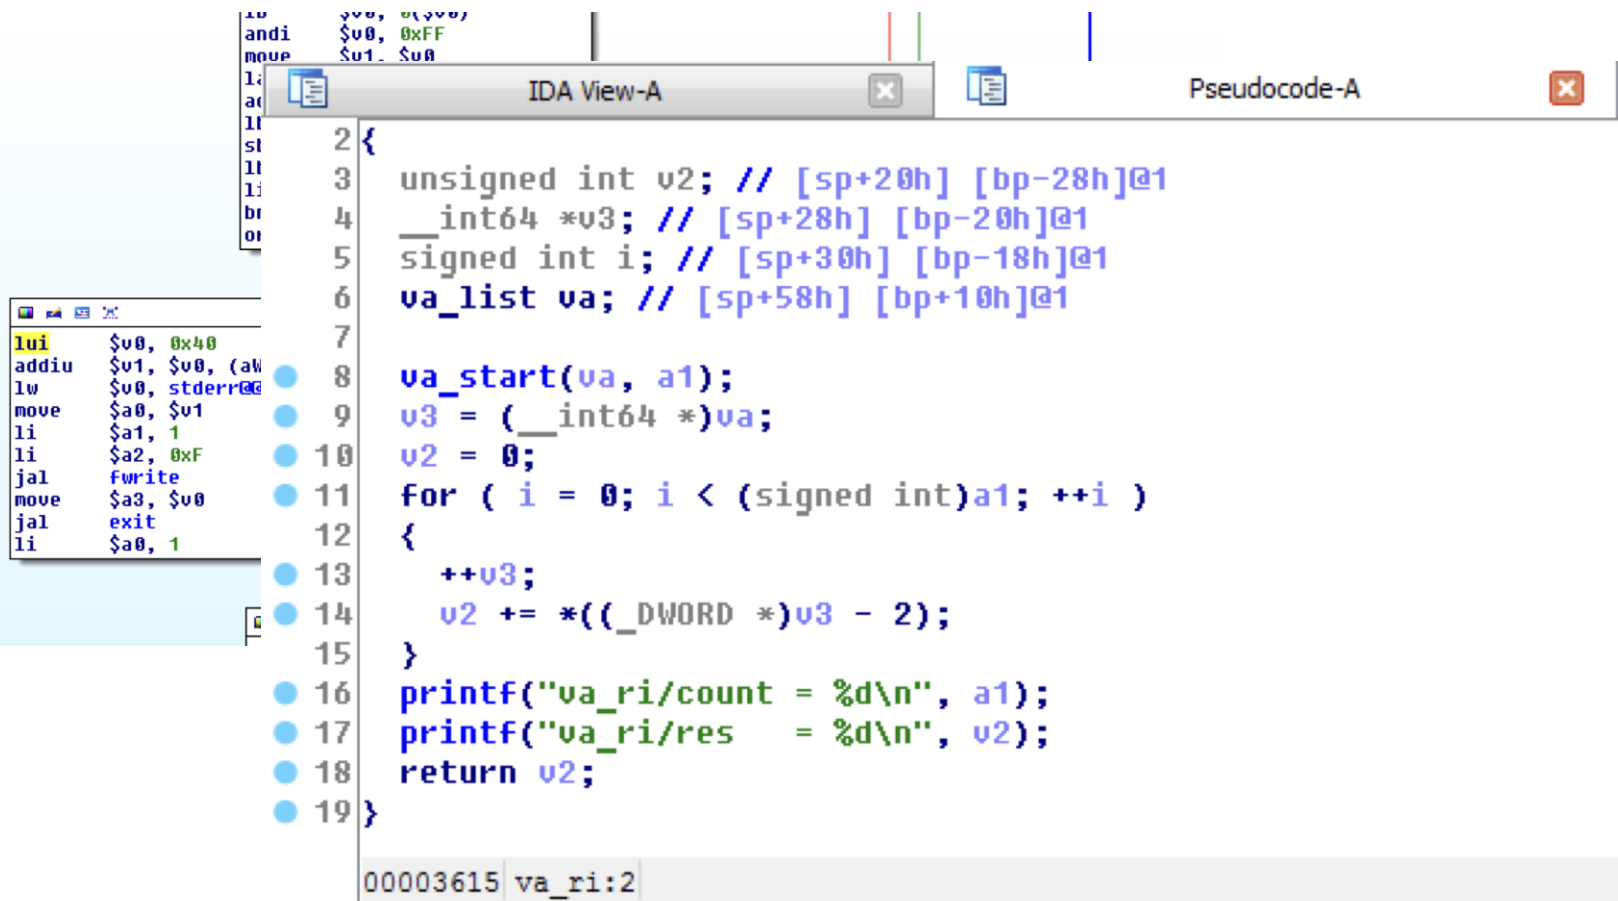


What if we start with binary?



Very complex, usually poor results

# Decompilation



The screenshot displays the IDA Pro interface with two windows open: 'IDA View-A' and 'Pseudocode-A'. The 'IDA View-A' window shows assembly code for a function, with the following instructions visible:

```
lui    $v0, 0x40
addiu  $v1, $v0, (a1
lw      $v0, stderr@G
move    $a0, $v1
li      $a1, 1
li      $a2, 0xF
jal      fwrite
move    $a3, $v0
jal      exit
li      $a0, 1
```

The 'Pseudocode-A' window shows the decompiled C++ code for the same function:

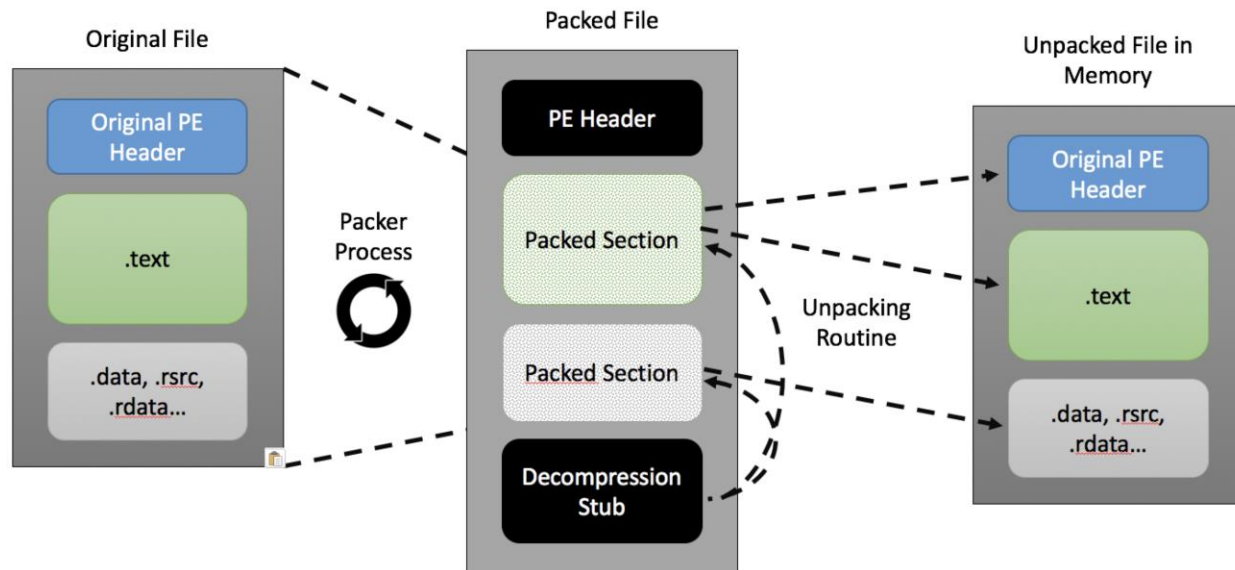
```
2 {
3     unsigned int v2; // [sp+20h] [bp-28h]@1
4     __int64 *v3; // [sp+28h] [bp-20h]@1
5     signed int i; // [sp+30h] [bp-18h]@1
6     va_list va; // [sp+58h] [bp+10h]@1
7
8     va_start(va, a1);
9     v3 = (__int64 *)va;
10    v2 = 0;
11    for ( i = 0; i < (signed int)a1; ++i )
12    {
13        ++v3;
14        v2 += *((_DWORD *)v3 - 2);
15    }
16    printf("va_ri/count = %d\n", a1);
17    printf("va_ri/res = %d\n", v2);
18    return v2;
19 }
```

At the bottom of the 'Pseudocode-A' window, the address 00003615 is shown next to the variable va\_ri:2.

# Packing

Packing hides the real code of a program through one or more layers of compression/encryption

At run-time the unpacking routine restores the original code in memory and then executes it





# Vulnerability Discovery

 [http://www.immunityinc.com/downloads/DaveAitel\\_TheHackerStrategy.pdf](http://www.immunityinc.com/downloads/DaveAitel_TheHackerStrategy.pdf)

Experienced analysts (according to Aitel)...

1 hour of binary analysis:

- Simple backdoors, coding style, bad API calls (strcpy)

1 week of binary analysis:

- Likely to find 1 good vulnerability

1 month of binary analysis:

- Likely to find 1 vulnerability no one else will ever find

# How to Find Vulnerabilities?

---

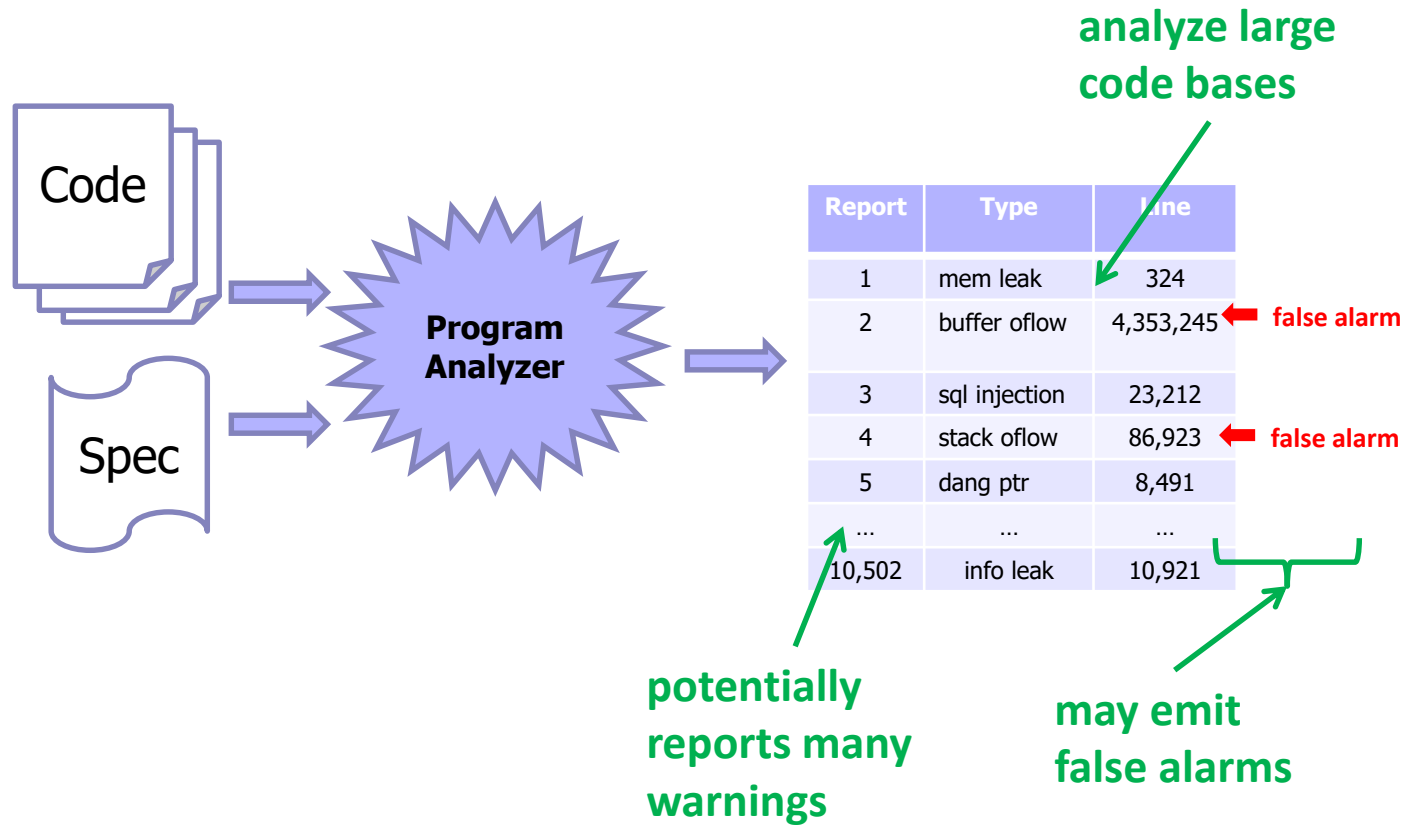
## Manual analysis

- Source code review
- Reverse engineering

## Program analysis tools:

- Static analysis
- Fuzzing
- Symbolic analysis

# Program Analyzers



# False Positives & False Negatives

---

Term	Definition
False positive	A spurious warning that does not indicate an actual vulnerability
False negative	Does not emit a warning for an actual vulnerability

**Complete** analysis: no false negatives

**Sound** analysis: no false positives

# Soundness and Completeness

	Complete	Incomplete
Sound	<p>Reports all errors Reports no false alarms</p> <p>No false positives No false negatives</p> <p><b>Undecidable</b></p>	<p>Reports all errors May report false alarms</p> <p>No false negatives False positives</p> <p><b>Decidable</b></p>
Unsound	<p>May not report all errors Reports no false alarms</p> <p>False positives No false negatives</p> <p><b>Decidable</b></p>	<p>May not report all errors May report false alarms</p> <p>False negatives False positives</p> <p><b>Decidable</b></p>

# Example Tools

---

Approach	Type	Comment
Lexical analyzers	Static analysis	Perform syntactic checks  Ex: LINT, RATS, ITS4
Fuzz testing	Dynamic analysis	Run on specially crafted inputs to test
Symbolic execution	Emulated execution	Run program on many inputs at once  Ex: KLEE, S2E, FiE
Model checking	Static analysis	Abstract program to a model, check that model satisfies security properties  Ex: MOPS, SLAM, etc.

# Source Code Scanners

---

Program that looks at source code, flags suspicious constructs

```
...  
strcpy( ptr1, ptr2 );  
...
```

Warning: Don't use strcpy

Simplest example: grep

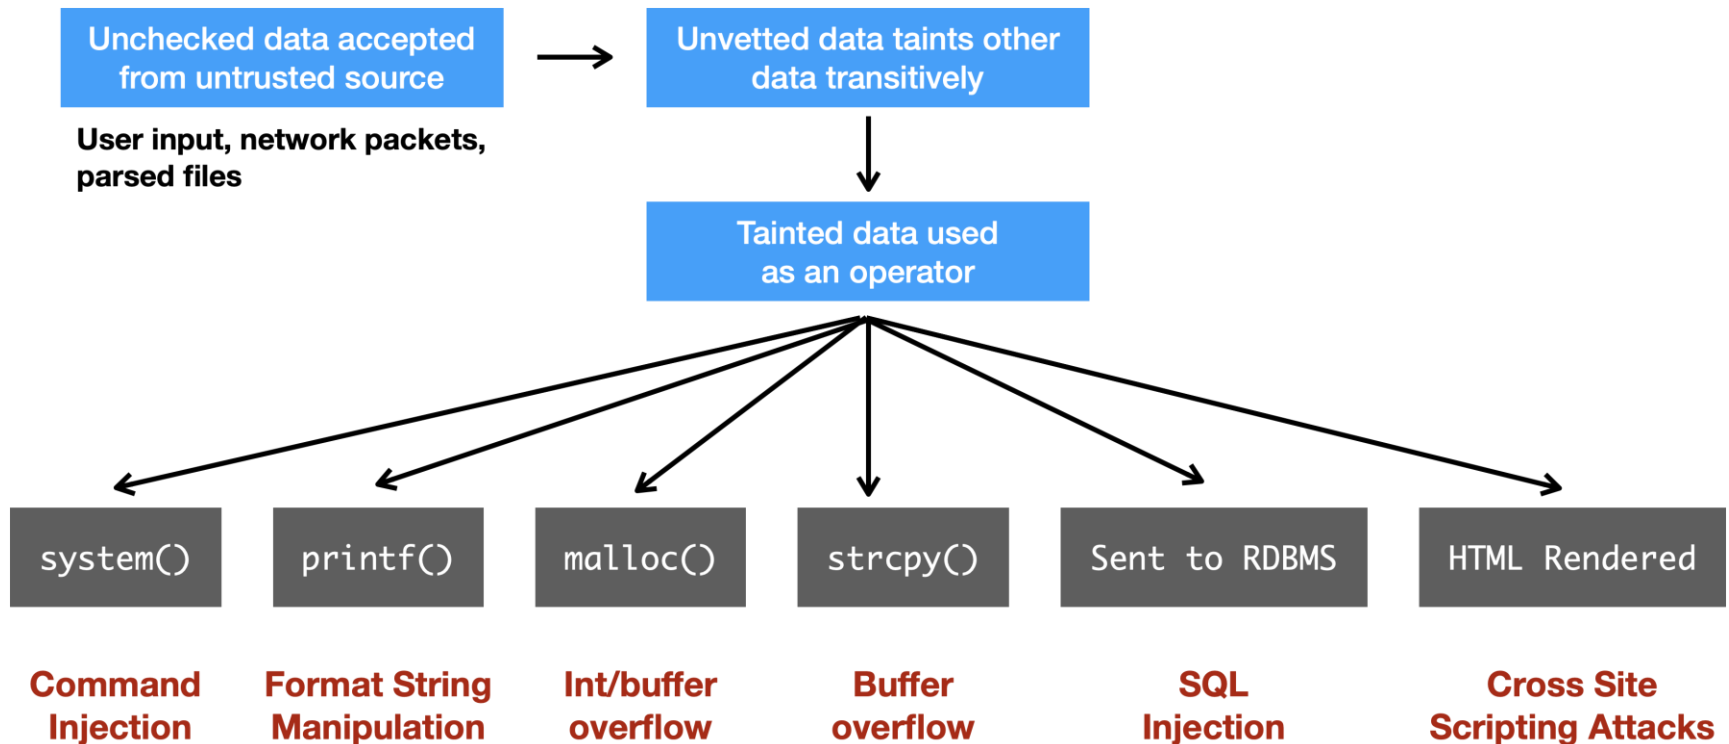
Lint is early example

RATS (Rough auditing tool for security)

ITS4 (It's the Software Stupid Security Scanner)

Circa 1990's technology, **shouldn't** work for reasonable modern codebases (... but probably will)

# Tainting Checkers





# Dynamic Analysis: Fuzzing

---

Choose a bunch of inputs  
See if they cause program to crash

Key challenge: **finding good inputs**




“The term first originates from a class project at the University of Wisconsin 1988 although similar techniques have been used in the field of quality assurance, where they are referred to as robustness testing, syntax testing or negative testing.”

[http://en.wikipedia.org/wiki/Fuzz\\_testing](http://en.wikipedia.org/wiki/Fuzz_testing)

# Fuzzing

---



# HTTP Fuzzing Example

---

## Standard HTTP GET request

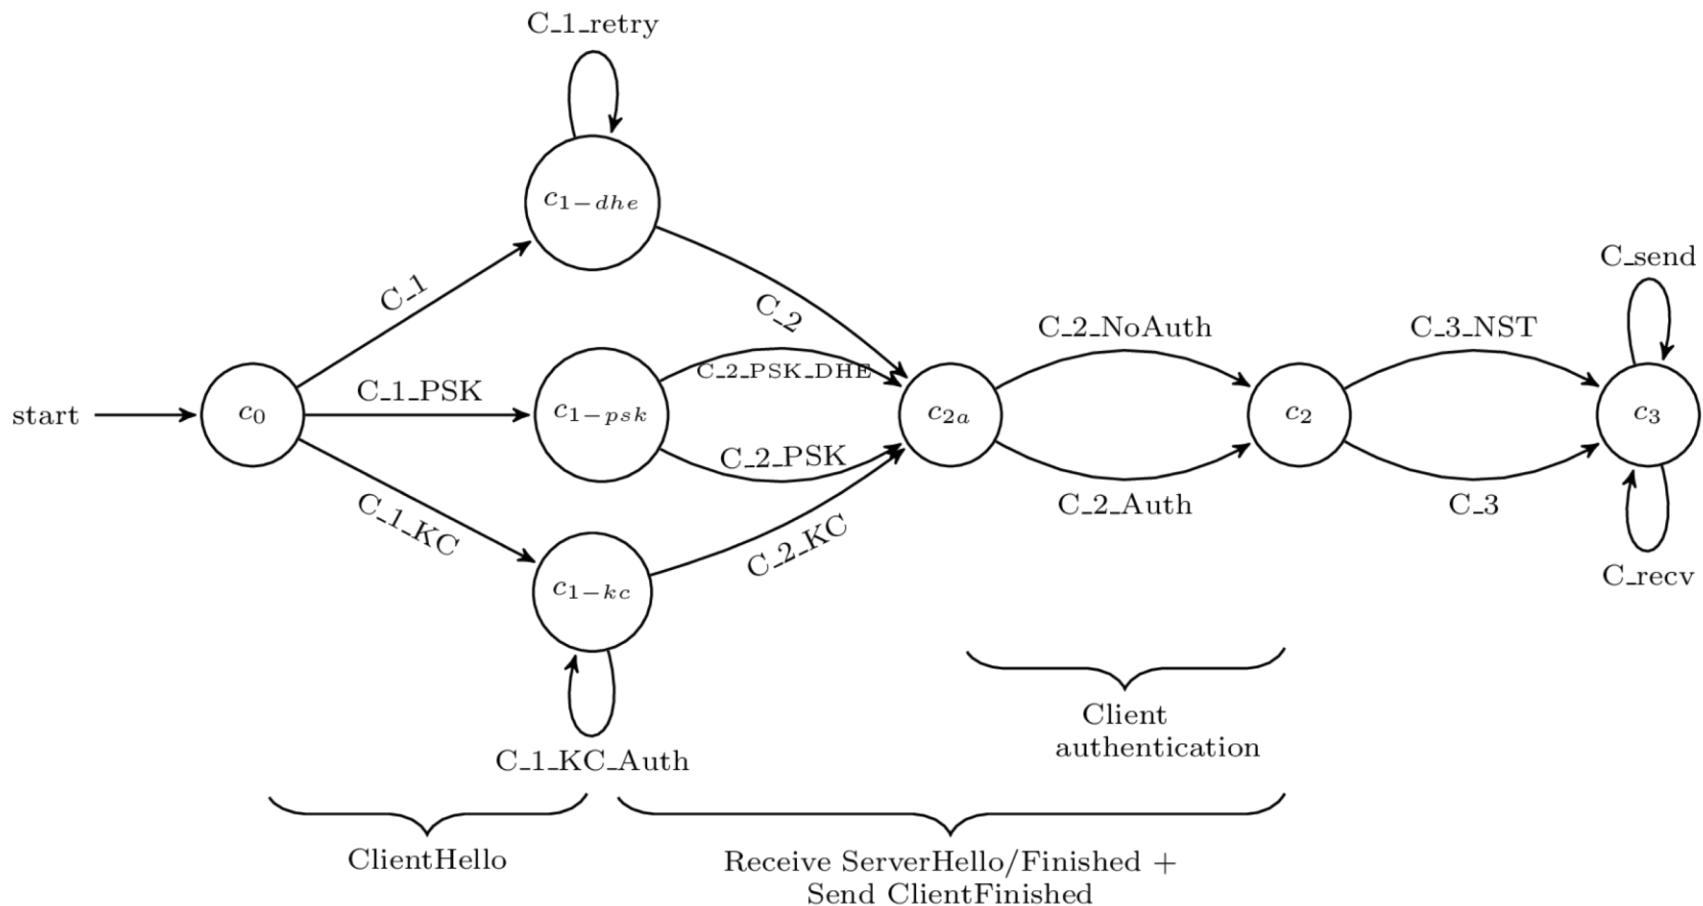
- GET /index.html HTTP/1.1

## Anomalous requests

- GEEEE...EET /index.html HTTP/1.1
- GET //////////index.html HTTP/1.1
- GET %n%n%n%n%n%n.html HTTP/1.1
- GET /AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA.html HTTP/1.1
- GET /index.html HTTPTTTTTTTTTTTTTTTP/1.1
- GET /index.html HTTP/1.1.1.1.1.1.1.1

but not df%w3rasd8#r78jskdflasdjf (why?)

# Problem with Random Fuzzing



TLS 1.3 state diagram

# Fuzzing

argv[1]="AAAA"  
argv[2]=1

Program

argv[1] = random str  
argv[2] = random 32-bit int

Program

If integers are 32 bits, then probability  
of crashing is **at most what?**  $1/2^{32}$

Achieving code coverage can  
be very difficult

```
main( int argc, char* argv[] ) {  
    char* b1;  
    char* b2;  
    char* b3;  
  
    if( argc != 3 ) then return 0;  
    if( atoi(argv[2]) != 31337 )  
        complicatedFunction();  
    else {  
        b1 = (char*)malloc(248);  
        b2 = (char*)malloc(248);  
        free(b1);  
        free(b2);  
        b3 = (char*)malloc(512);  
        strncpy( b3, argv[1], 511 );  
        free(b2);  
        free(b3);  
    }  
}
```

# Code Coverage and Fuzzing

---

Code coverage defined in many ways

- # of basic blocks reached
- # of paths followed
- # of conditionals followed
- gcov is useful standard tool

Mutation-based

- Start with known-good examples, mutate them to new cases
  - heuristics: increase string lengths (AAAAAAAAAA...)
  - randomly change items

Generative

- Start with specification of protocol, file format
- Build test case files from it
  - Rarely used parts of spec

# Generation Example

```
1  <!-- A. Local file header -->
2  <Block name="LocalFileHeader">
3    <String name="lfh_Signature" valueType="hex" value="504b0304" token="true" mut
4    <Number name="lfh_Ver" size="16" endian="little" signed="false"/>
5    ...
6    [truncated for space]
7    ...
8    <Number name="lfh_CompSize" size="32" endian="little" signed="false">
9      <Relation type="size" of="lfh_CompData"/>
10   </Number>
11   <Number name="lfh_DecompSize" size="32" endian="little" signed="false"/>
12   <Number name="lfh_FileNameLen" size="16" endian="little" signed="false">
13     <Relation type="size" of="lfh_FileName"/>
14   </Number>
15   <Number name="lfh_ExtraFldLen" size="16" endian="little" signed="false">
16     <Relation type="size" of="lfh_FldName"/>
17   </Number>
18   <String name="lfh_FileName"/>
19   <String name="lfh_FldName"/>
20   <!-- B. File data -->
21   <Blob name="lfh_CompData"/>
22 </Block>
```

# Mutation vs. Generation

---

	Ease of Use	Knowledge	Completeness	Complex Programs
Mutation	Easy to setup and automate	Little to no protocol knowledge required	Limited by initial corpus	May fail for protocols with checksums or other complexity
Generative	Writing generator is labor intensive	Requires having protocol specification	More complete than mutations	Handles arbitrarily complex protocols



# Evolutionary Fuzzing

---

Generate inputs based on the structure and **response** of the program

Autodafe: Prioritizes based on inputs that reach dangerous API functions

EFS: Generates test cases based on code coverage metrics

Typically instrument program with additional instructions to track what code has been reached — or, if no source is available, track with Valgrind.

# American Fuzzy Lop (AFL)


Widely used, highly effective fuzzing tool

- Specify example inputs
- Compile program with special afl compiler
- Run it

Performs mutation-based fuzzing:

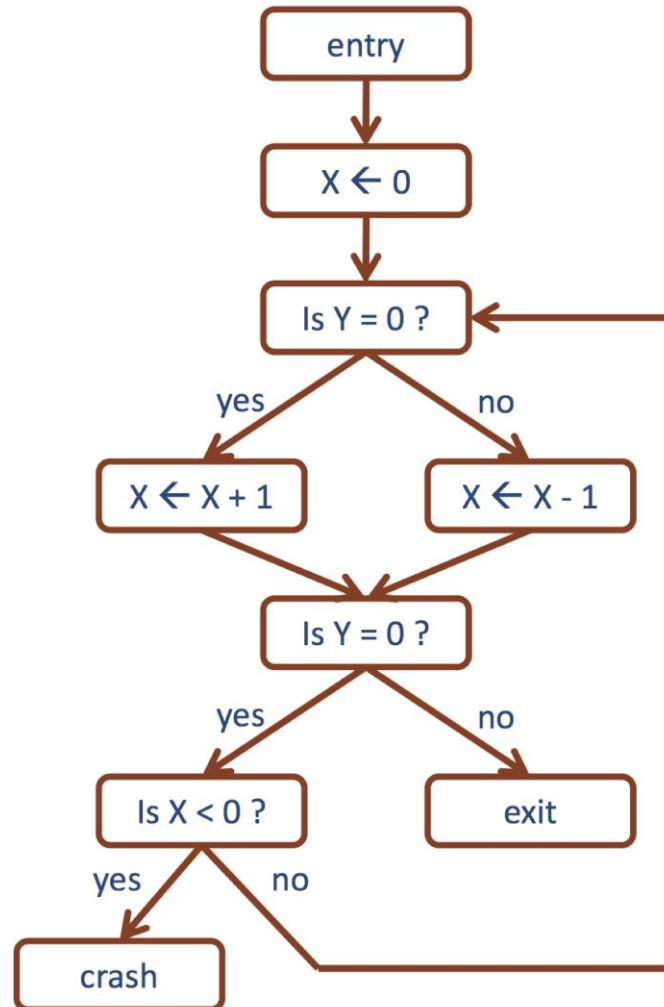
- Deterministic transforms to input (flip each bit, “walking byte flips”, etc.)
- Randomized stacked transforms
- Measure (approximation of) path coverage, keep and mutate set of files that increase coverage

Really fast and simple. Used to find bugs in Firefox, OpenSSH, BIND, ImageMagick, iOS, ...

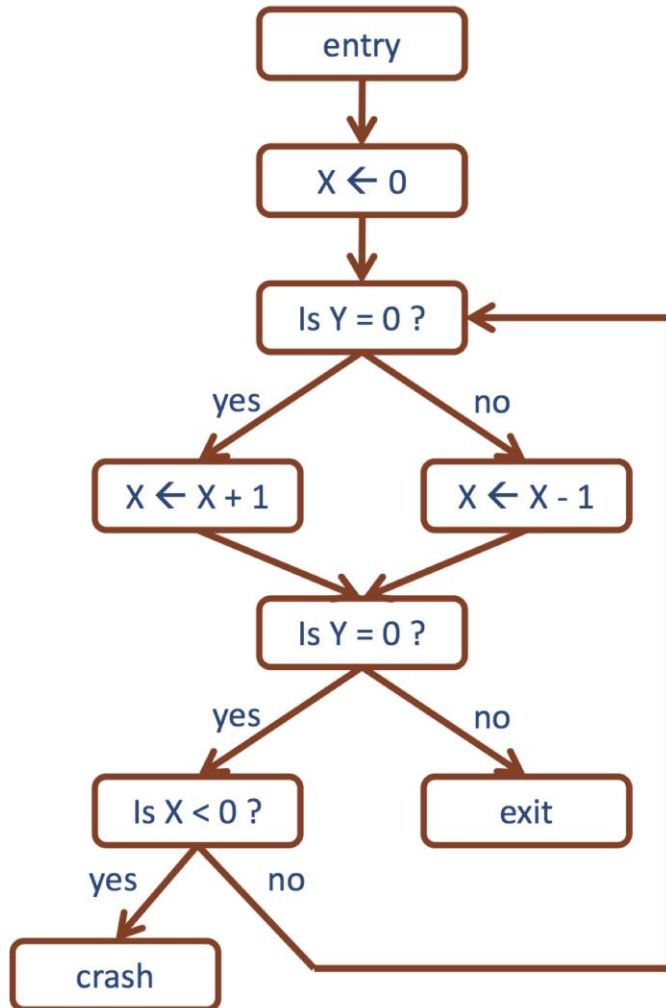


<https://lcamtuf.blogspot.com/2014/11/pulling-jpegs-out-of-thin-air.html>

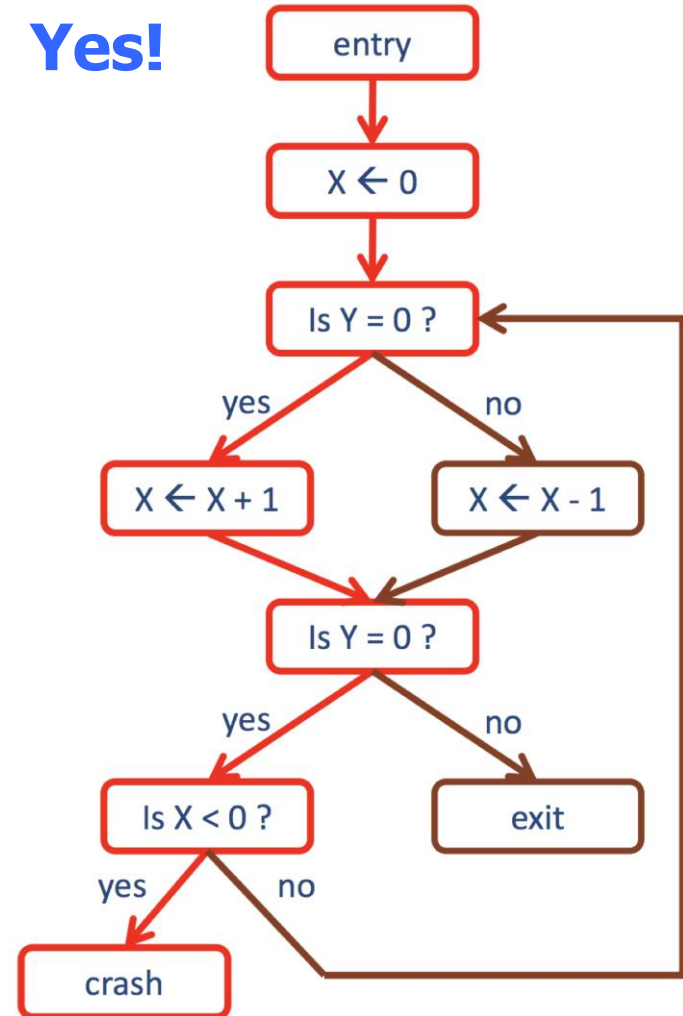
# Is This Program Safe?



# Is This Program Safe?

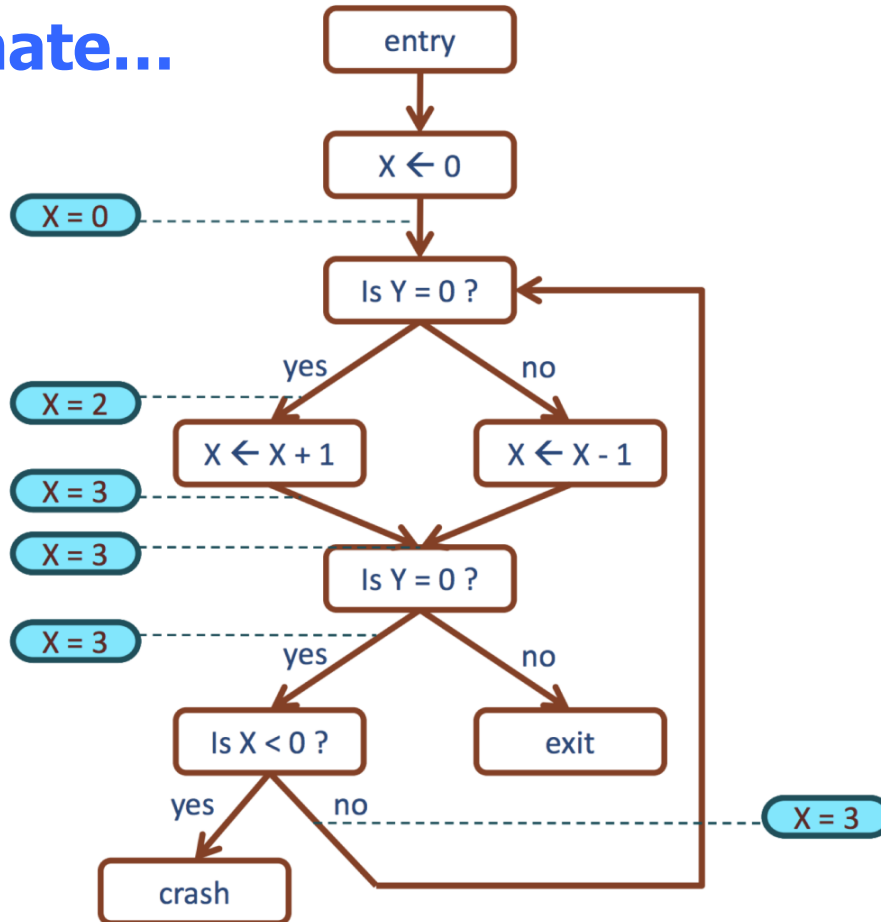


Yes!

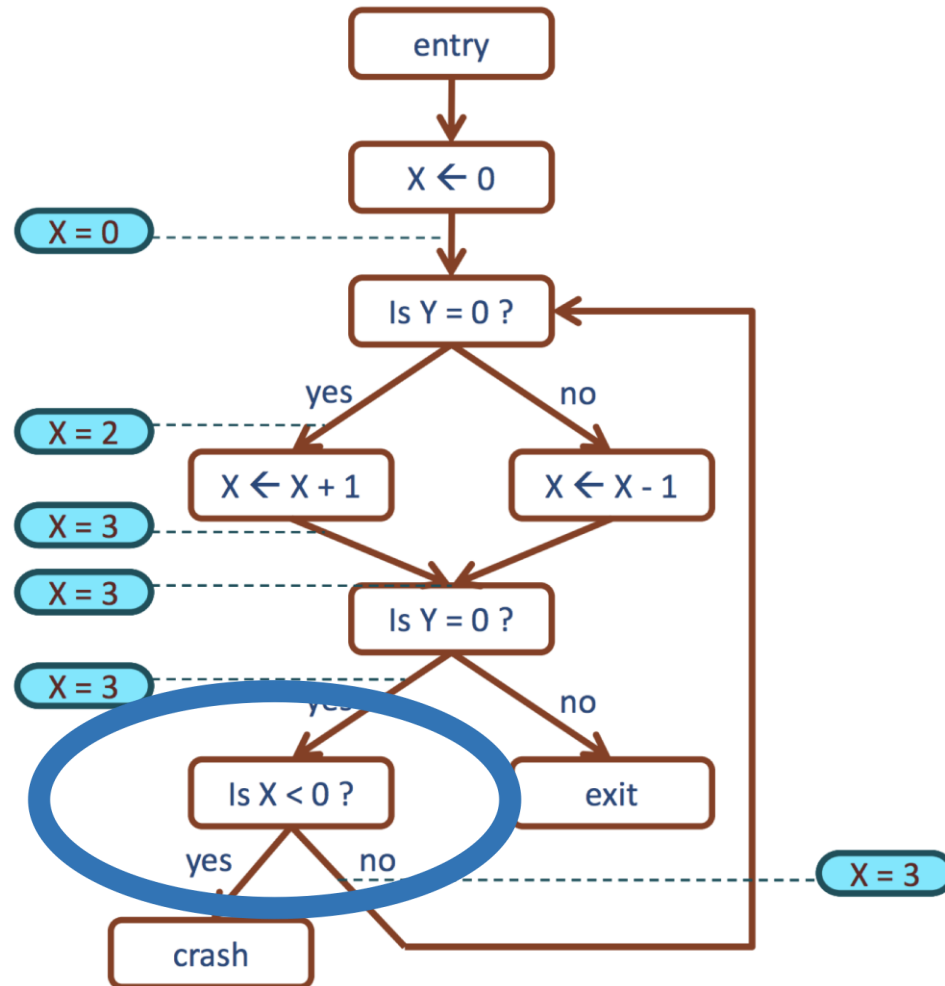


# Without Approximation...

Does not terminate...



# Abstract from Concrete Values



# Abstraction

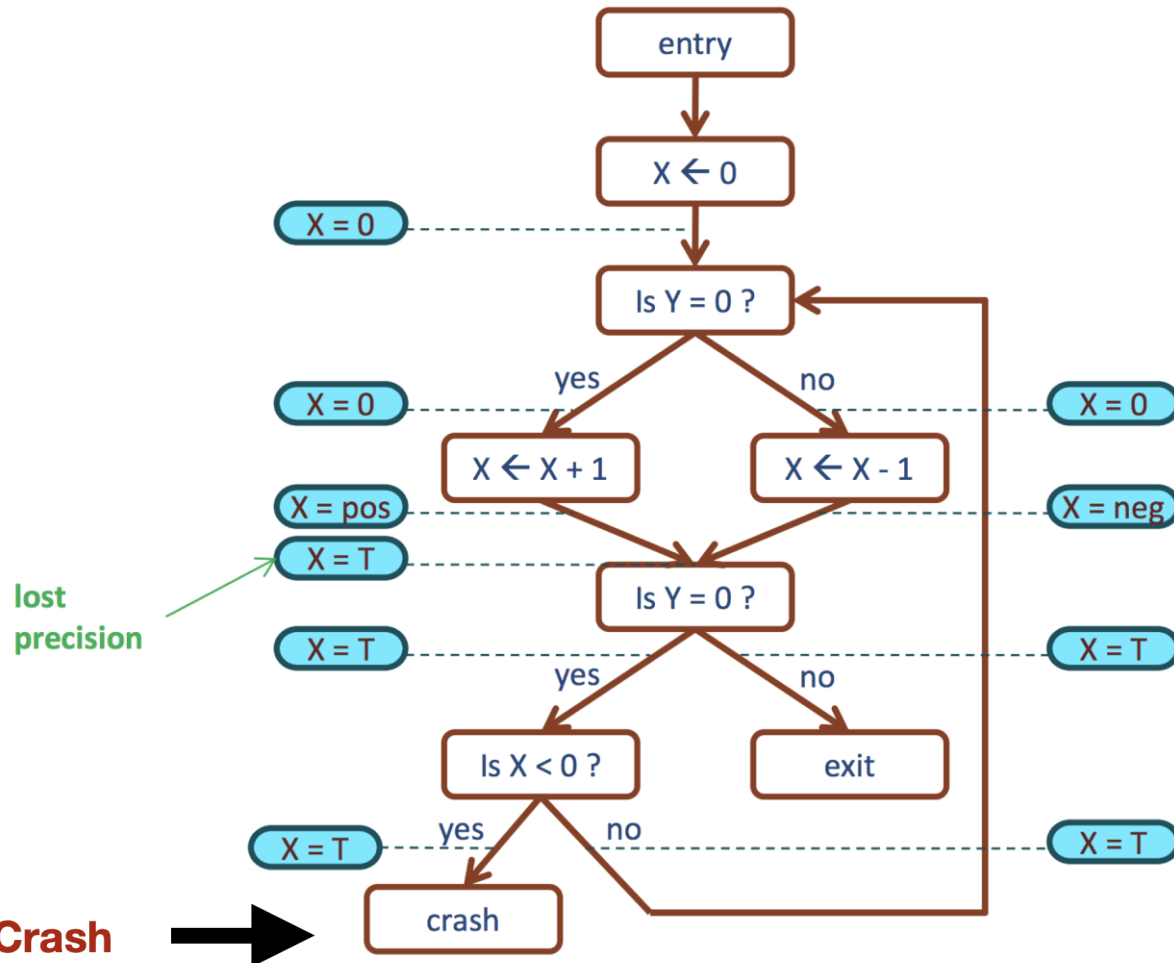
---

Concrete domain of integers

Abstract domain of signs

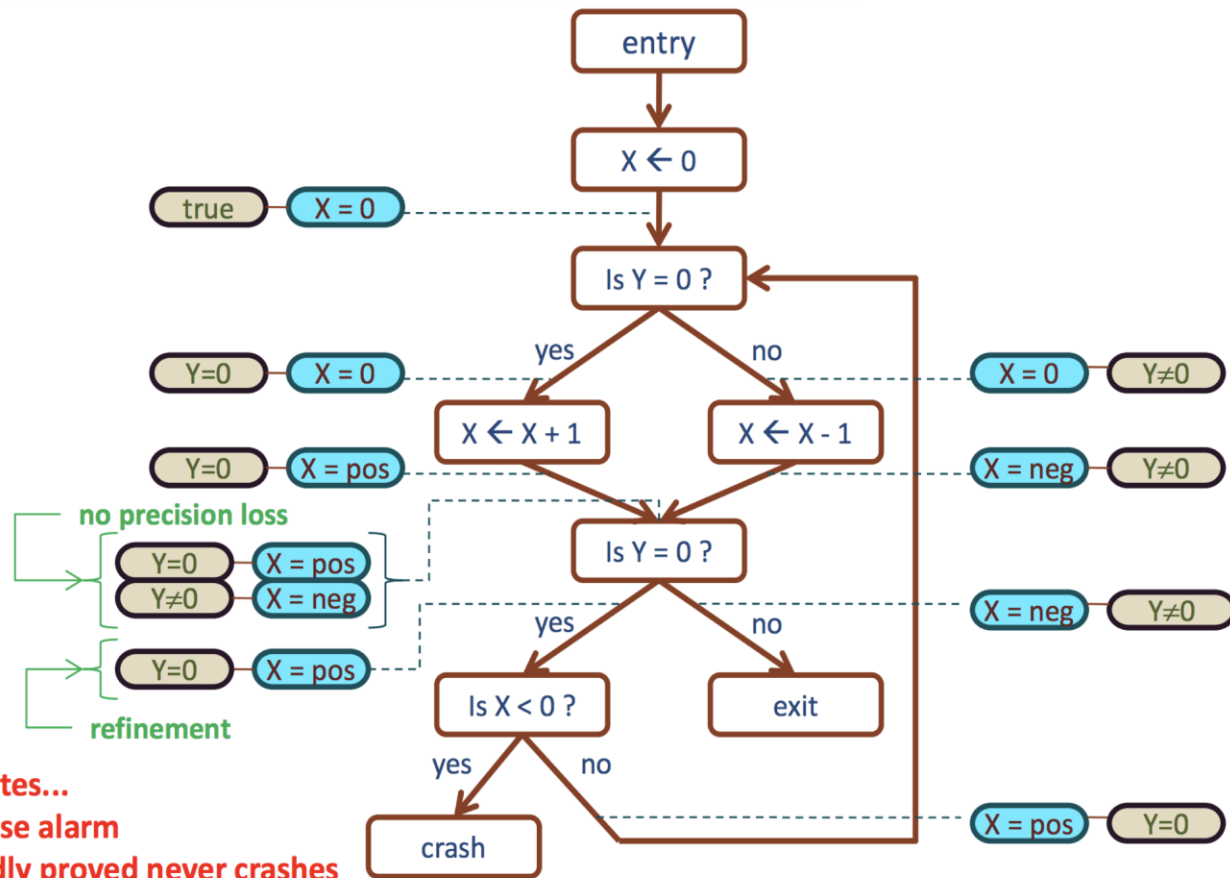
$x=5$	→	Positive integers
$x=-5$	→	Negative integers
$x=0$	→	Zero
$x=b \text{ ? } -1 : 1$	→	Integers
$x=y / 0$	→	No integers (undefined)

# With "Signs" Approximation



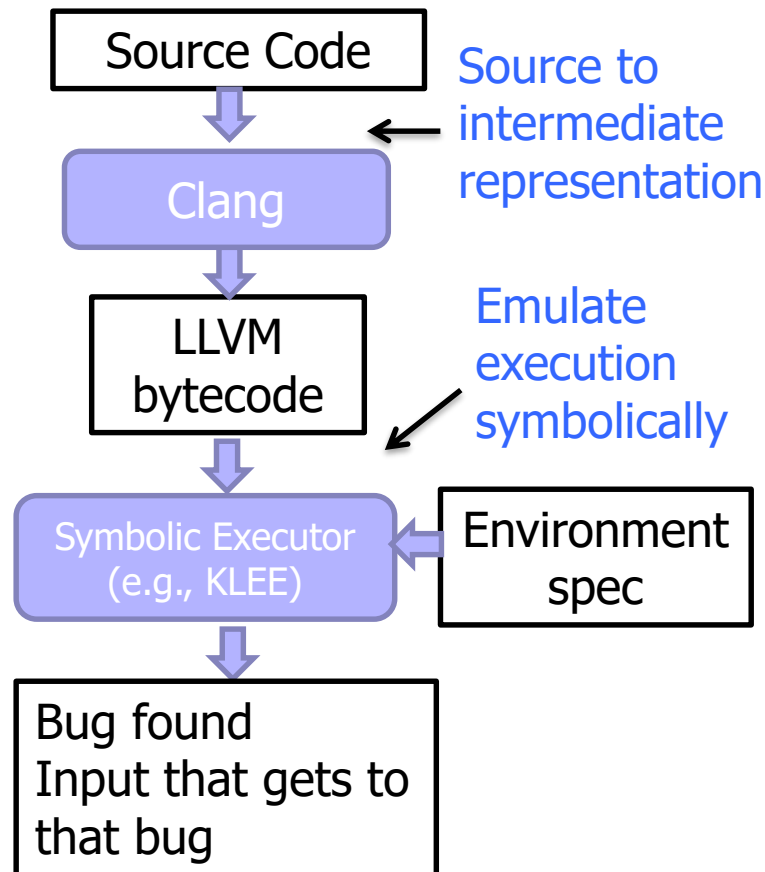


# Add Path Sensitivity



# Symbolic Execution

---



Technique for analyzing code paths and finding inputs

Associate **symbols** to input variables ("symbolic variable")

Simulate execution symbolically

- Update symbolic variable's value appropriately
- Conditionals add constraints on possible values

Cast constraints as satisfiability, use SAT solver to find inputs

Perform security checks at each execution state

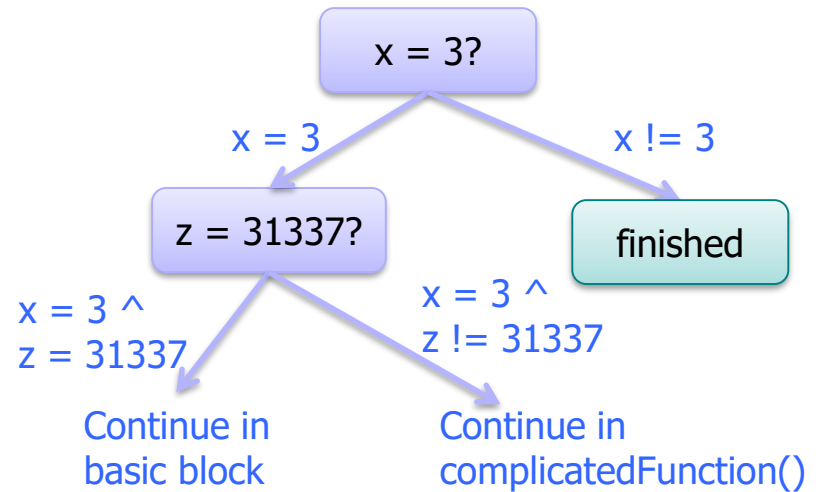
# Symbolic Execution

```
main( int argc, char* argv[] ) {  
    char* b1;  
    char* b2;  
    char* b3;  
  
    if( argc != 3 ) then return 0;  
    if( argv[2] != 31337 )  
        complicatedFunction();  
    else {  
        b1 = (char*)malloc(248);  
        b2 = (char*)malloc(248);  
        free(b1);  
        free(b2);  
        b3 = (char*)malloc(512);  
        strncpy( b3, argv[1], 511 );  
        free(b2);  
        free(b3);  
    }  
}
```

Initially:

$argc = x$  (unconstrained int)

$argv[2] = z$  (memory array)



- Eventually emulation hits a double free
- Can trace back up path to determine what  $x, z$  must have been to hit this basic block

# Symbolic Execution Challenges

---

Can we complete analyses?

- Yes, but only for very simple programs
- Exponential # of paths to explore
- Each branch increases state size of symbolic emulator

Path selection

- Which state to explore next?
- Might get stuck in complicatedFunction()

Encoding checks on symbolic states

- Must include logic for double free check
- Symbolic execution on binary more challenging (lose most memory semantics)

# Example Tools

---

Approach	Type	Comment
Lexical analyzers	Static analysis	Perform syntactic checks  Ex: LINT, RATS, ITS4
Fuzz testing	Dynamic analysis	Run on specially crafted inputs to test
Symbolic execution	Emulated execution	Run program on many inputs at once, by  Ex: KLEE, S2E, FiE
Model checking	Static analysis	Abstract program to a model, check that model satisfies security properties  Ex: MOPS, SLAM, etc.

# Google Address Sanitizer (ASan)

---

Memory error detector for C/C++ that finds...

- Use after free (dangling pointer dereference)
- Heap buffer overflow
- Stack buffer overflow
- Global buffer overflow
- Use after return
- Use after scope
- Initialization order bugs
- Memory leaks

# Google Address Sanitizer (ASan)

---

## LLVM Pass

- Modifies the code to check the shadow state for each memory access and creates poisoned redzones around stack and global objects to detect overflows and underflows

A run-time library that replaces memory management functions

- Replaces malloc, free and related functions, creates poisoned redzones around allocated heap regions, delays the reuse of freed heap regions, and does error reporting

# Google Address Sanitizer (ASan)

---

==9901==ERROR: AddressSanitizer: heap-use-after-free on address 0x60700000dfb5 at pc 0x45917b bp 0x7fff4490c700 sp 0x7fff4490c6f8

READ of size 1 at 0x60700000dfb5 thread T0

#0 0x45917a in main use-after-free.c:5

#1 0x7fce9f25e76c in \_\_libc\_start\_main /build/builddd/eglibc-2.15/csu/libc-start.c:226

#2 0x459074 in \_start (a.out+0x459074)

0x60700000dfb5 is located 5 bytes inside of 80-byte region [0x60700000dfb0,0x60700000e000) freed by thread T0 here:

#0 0x4441ee in \_\_interceptor\_free projects/compiler-rt/lib/asan/asan\_malloc\_linux.cc:64

#1 0x45914a in main use-after-free.c:4

#2 0x7fce9f25e76c in \_\_libc\_start\_main /build/builddd/eglibc-2.15/csu/libc-start.c:226

previously allocated by thread T0 here:

#0 0x44436e in \_\_interceptor\_malloc projects/compiler-rt/lib/asan/asan\_malloc\_linux.cc:74

#1 0x45913f in main use-after-free.c:3

#2 0x7fce9f25e76c in \_\_libc\_start\_main /build/builddd/eglibc-2.15/csu/libc-start.c:226

**SUMMARY: AddressSanitizer: heap-use-after-free use-after-free.c:5 main**



# Summary of Program Analysis

---

	Pros	Cons
Static	Enables quickly finding bugs at development time Can detect some problems that dynamic misses	Either over or under reports. Misses complex bugs. Generally requires code.
Dynamic	May uncover complex behavior missed by static. Can run on blackbox.	Depends on user input— only checks executed code

# Bug Finding is a Big Business

---

Grammatech (Cornell startup, 1988)

Coverity (Stanford startup)

- Great article on static analysis in the real world:  
<http://web.stanford.edu/~engler/BLOC-coverity.pdf>

Fortify

... many, many others

Also reverse engineers, exploit developers,  
zero-day markets...