

Microarchitectural Attacks

Vitaly Shmatikov

Performance in Modern CPUs

Clock speed maxed out:

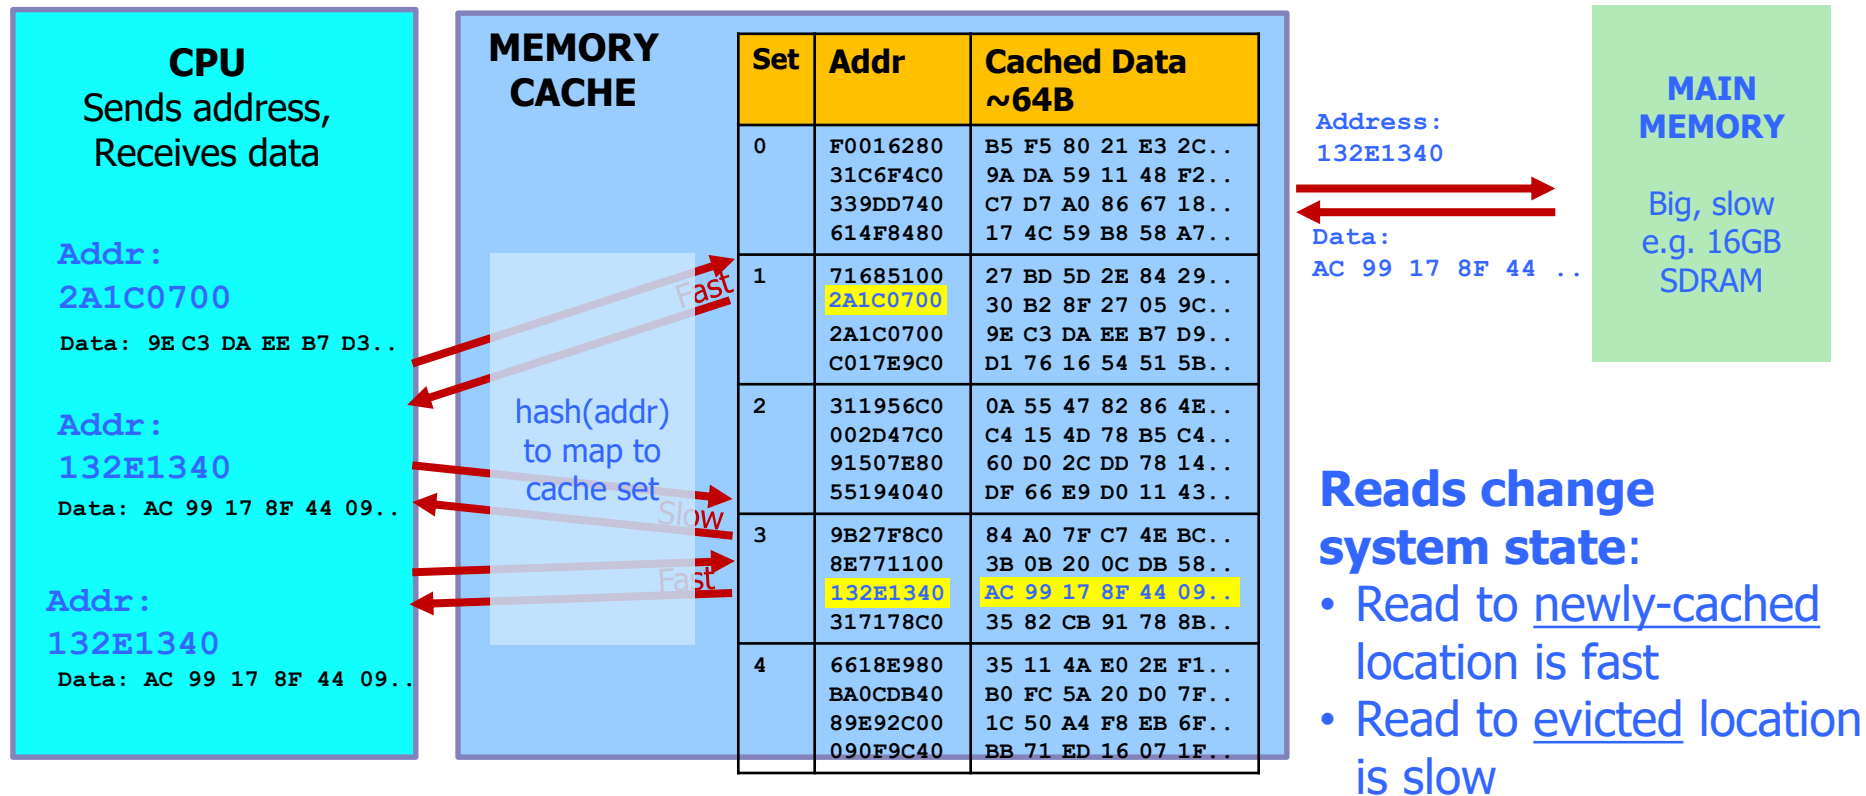
- Pentium 4 reached 3.8 GHz in 2004
- Memory latency is slow and not improving much

To gain performance, need to do more per cycle!


- Reduce memory delays: **caches**
- Work during delays: **speculative execution**

Memory Caches

Caches hold local (fast) copy of recently accessed 64-byte chunks of memory



Cross-VM Side Channel



Target is 4096-bit ElGamal secret key e

Modular Exponentiation (x, e, N):

let $e_n \dots e_1$ be the bits of e

$y \leftarrow 1$

for e_i in $\{e_n \dots e_1\}$

$y \leftarrow \text{Square}(y)$ (S)

$y \leftarrow \text{Reduce}(y, N)$ (R)

if $e_i = 1$ then

$y \leftarrow \text{Multi}(y, x)$ (M)

$y \leftarrow \text{Reduce}(y, N)$ (R)


return y // $y = x^e \bmod N$

$e_i = 1 \rightarrow \text{"SRMR"}$

$e_i = 0 \rightarrow \text{"SR"}$

Sequence of function calls
reveals secret key

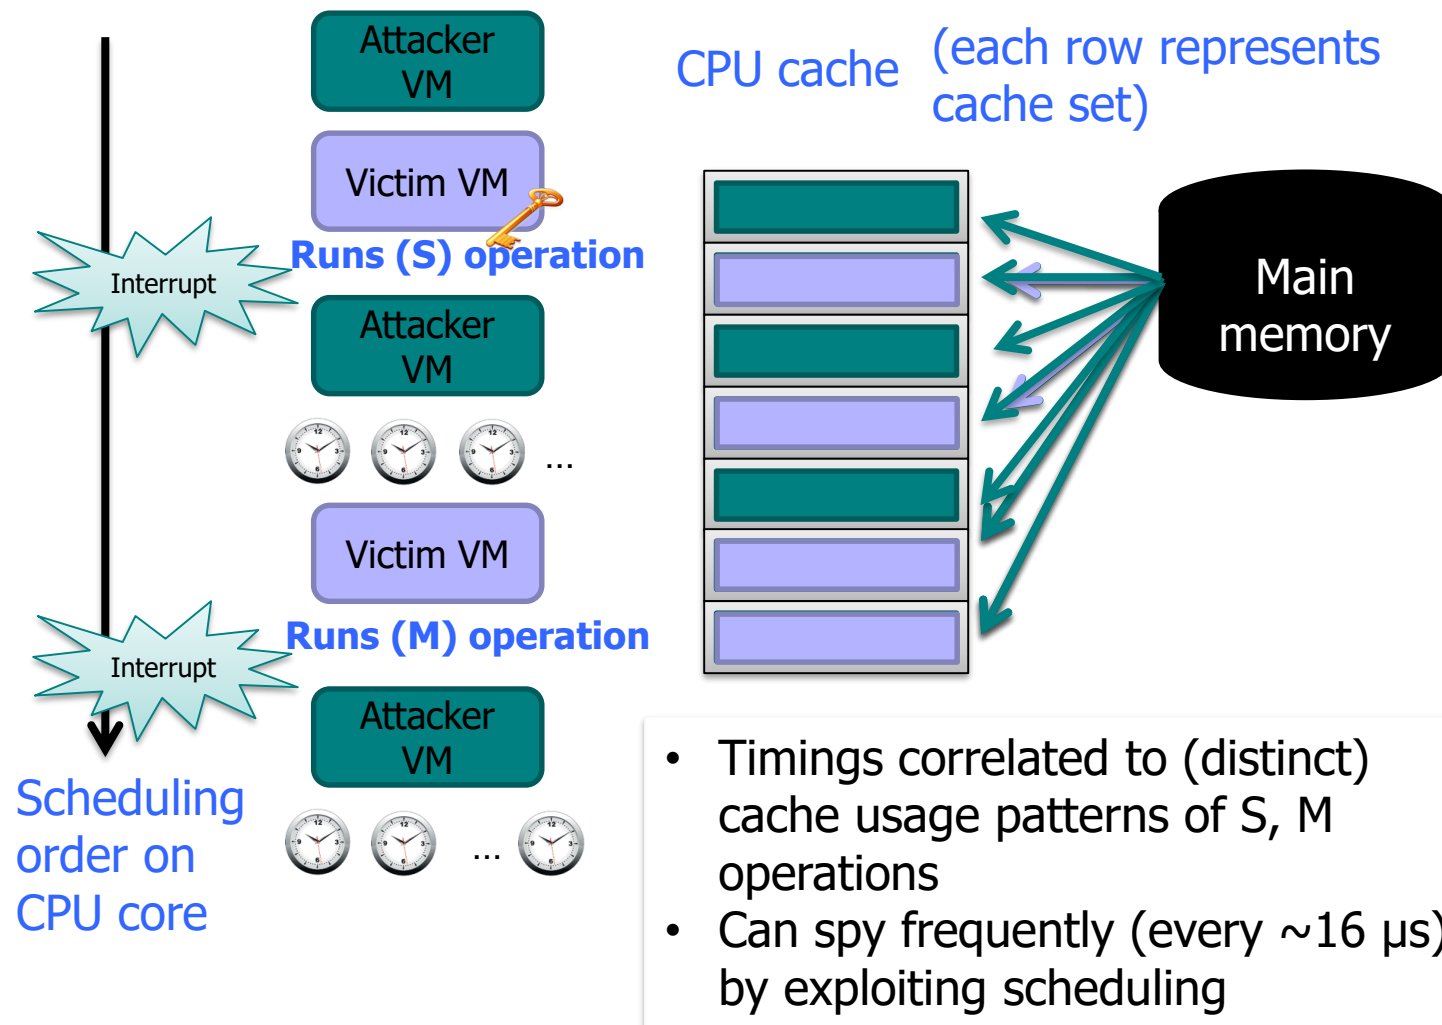
Cache Contention



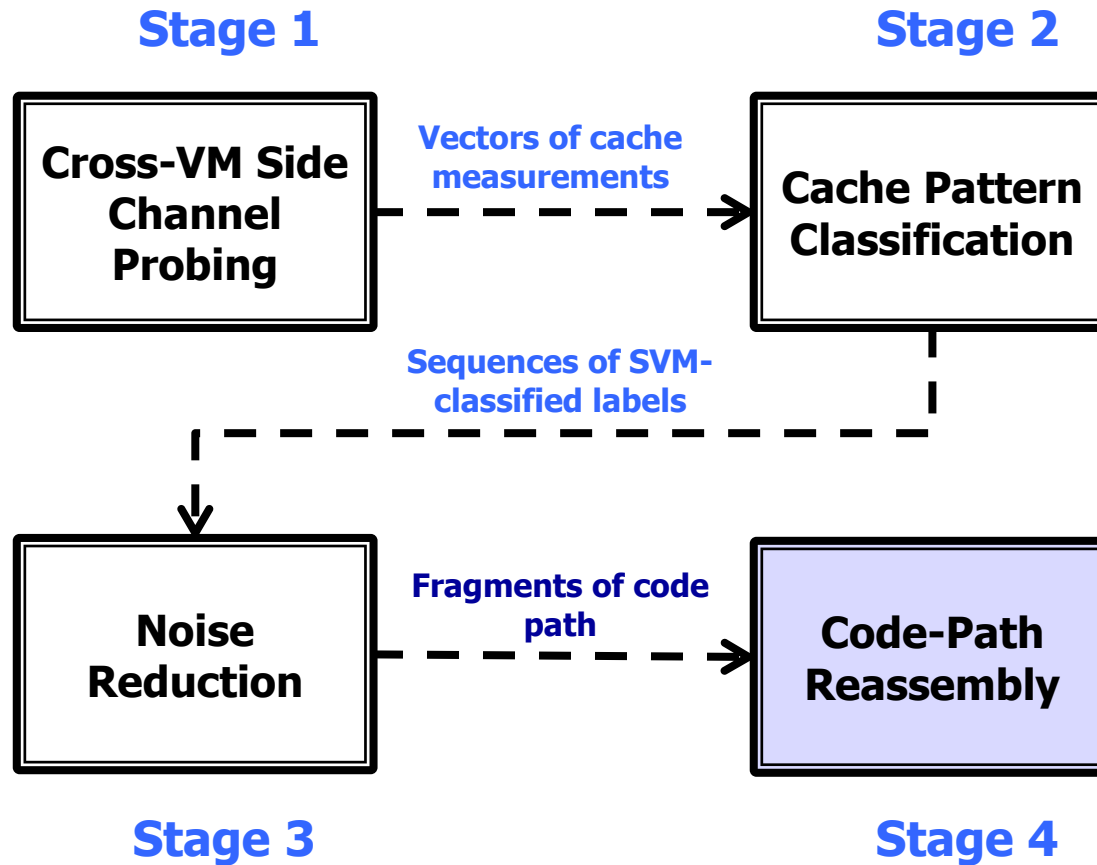
- 1) Read in a large array (fill CPU cache with attacker data)
- 2) Busy loop (allow victim to run)
- 3) Measure time to read large array (the load measurement)

Locations in cache occupied by victim will take longer to load ➡ Information about victim's use of cache revealed to attacker

Prime + Probe



Attack Stages



Prime + Probe Feasibility

Setup for in-lab experimentation:

- Intel Yorkfield processor (4 cores, 32KB L1 instruction cache)
- Xen + Linux + GnuPG + libgcrypt

Best result:

- 300,000,000 prime-probe results (6 hours)
- Over 300 key fragments
- Brute force the secret key in ~ 9800 guesses

Not practical in deployment settings

Microarchitectural Side Channels

Lots of research

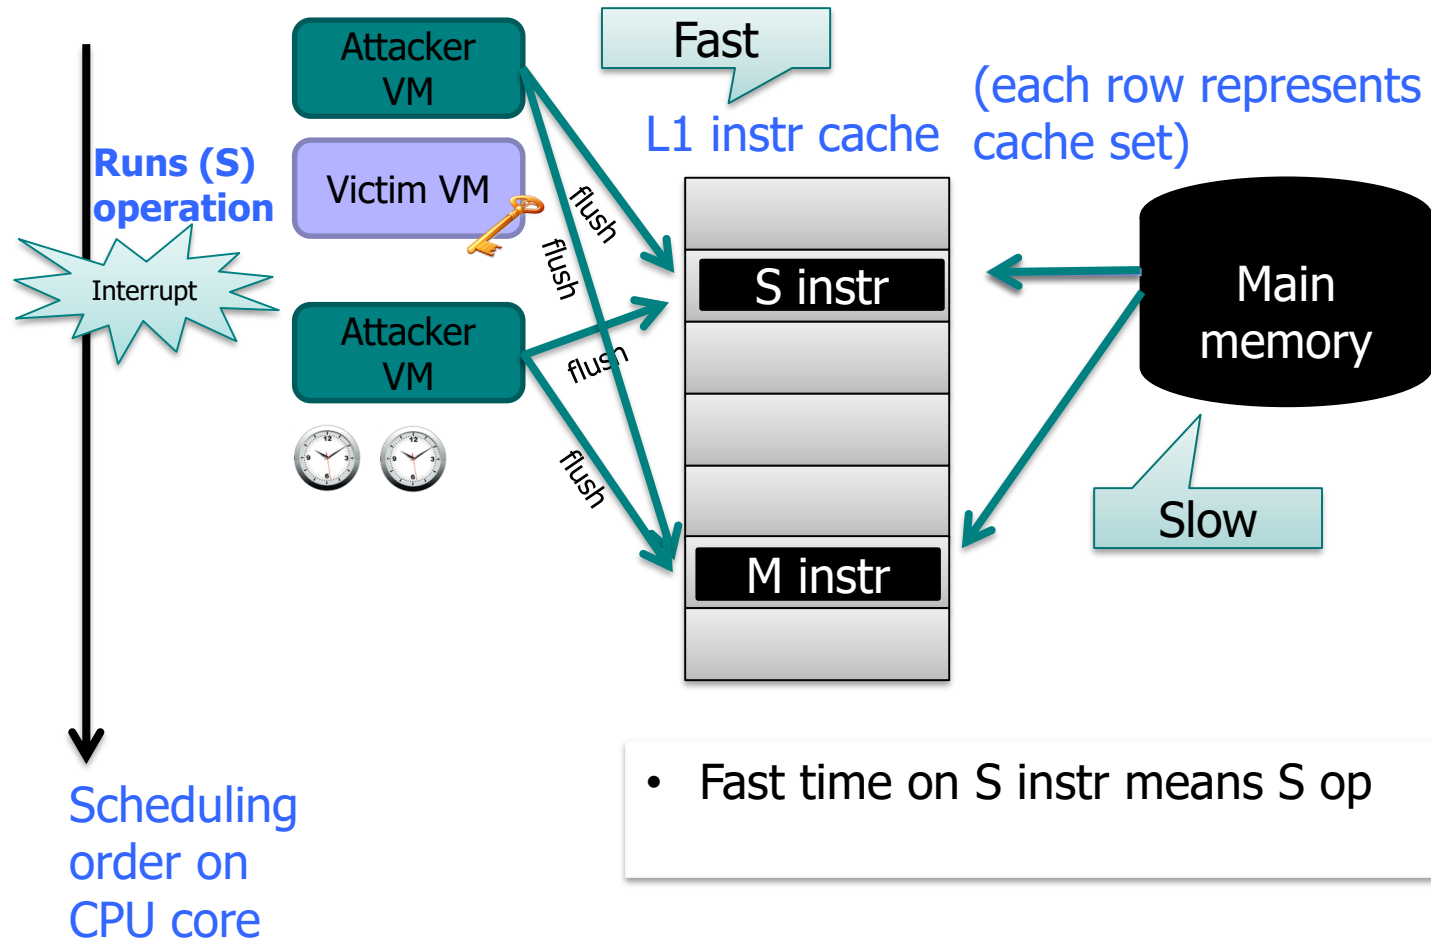
State-of-the-art Prime+Probe attacks

- Sinan Inci et al. 2016 "Cache Attacks Enable Bulk Key Recovery on the Cloud"

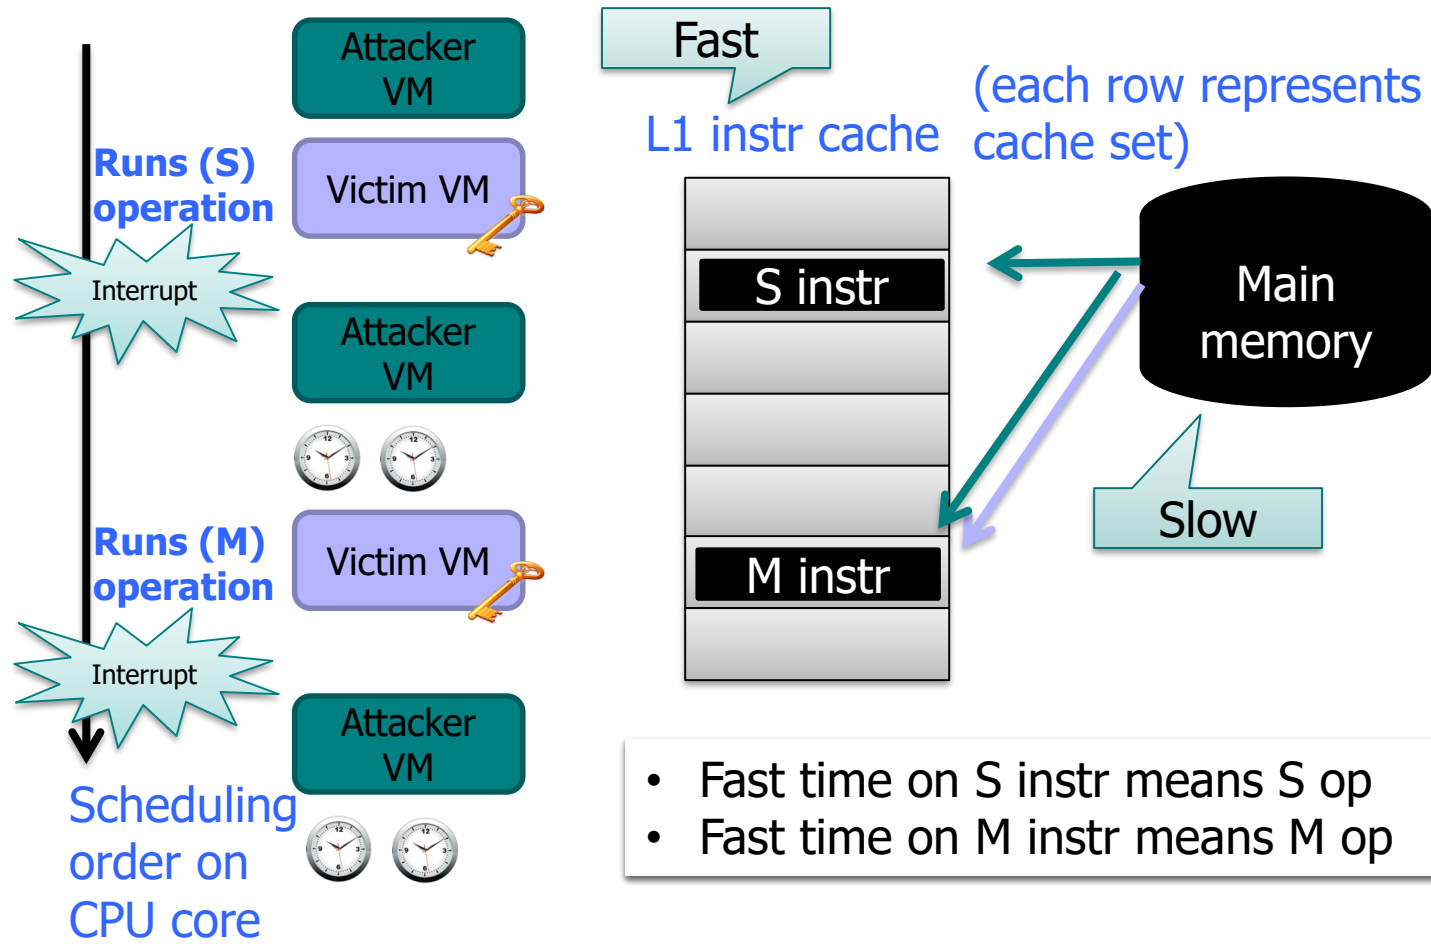
Flush+Reload more robust side channel in shared memory settings [Yarom, Falkner 2013]

- Spy process flushes memory shared with victim from caches
- Idles
- Times how long it takes to read shared memory

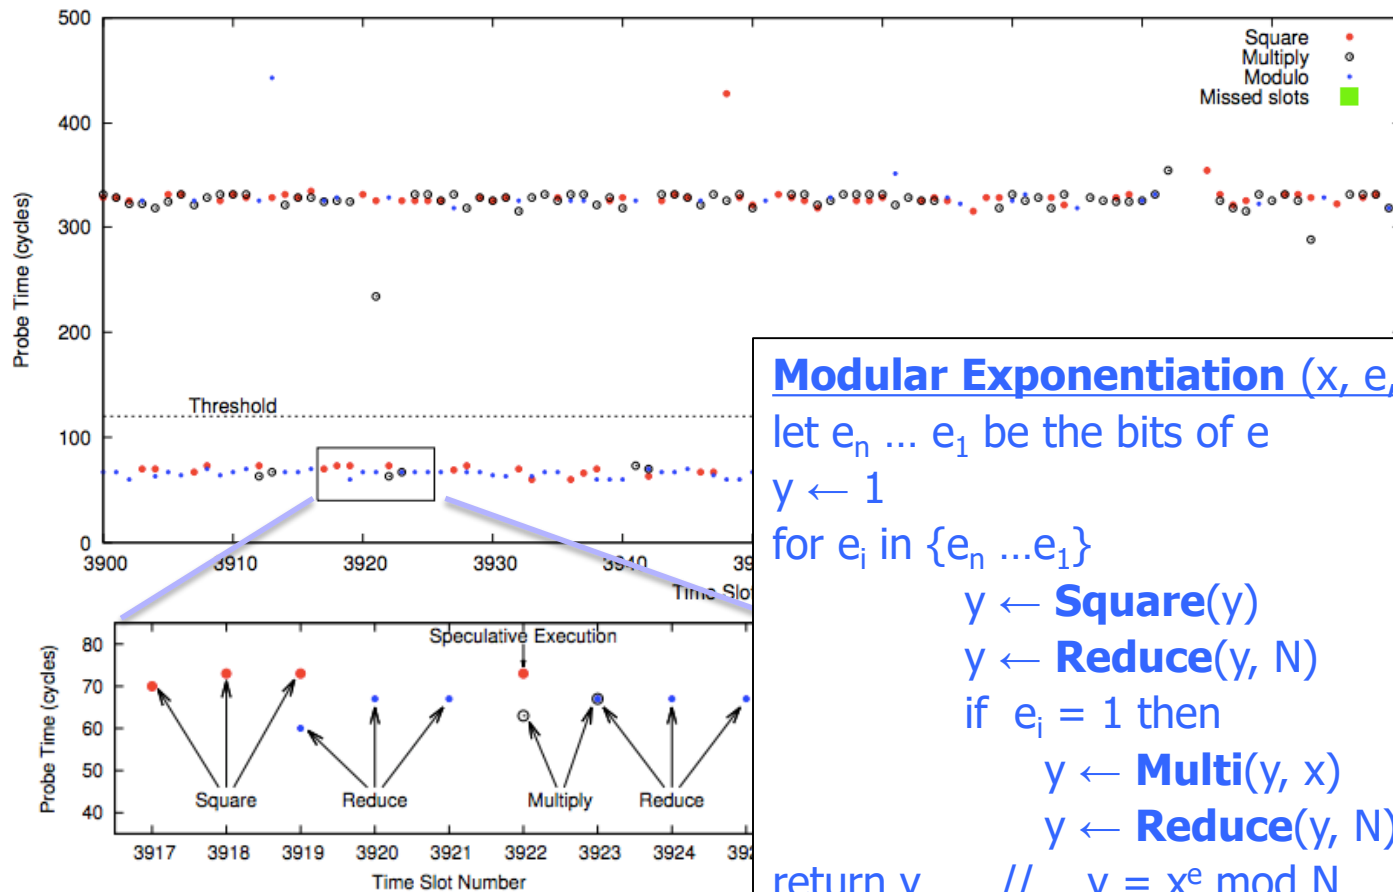
Flush + Reload



Flush + Reload



Attacking Square-and-Multiply



Modular Exponentiation (x, e, N) :

let $e_n \dots e_1$ be the bits of e

$y \leftarrow 1$

for e_i in $\{e_n \dots e_1\}$

$y \leftarrow \mathbf{Square}(y)$ **(S)**

$y \leftarrow \mathbf{Reduce}(y, N)$ **(R)**

if $e_i = 1$ then

$y \leftarrow \mathbf{Multi}(y, x)$ **(M)**

$y \leftarrow \mathbf{Reduce}(y, N)$ **(R)**

return y // $y = x^e \bmod N$

Speculative Execution

CPUs can guess likely program path and do speculative execution

‣ Example:

```
if (uncached_value == 1)    // load from memory
    a = compute(b)
```

- Branch predictor guesses if() is 'true' (based on prior history)
- Starts executing compute(b) speculatively
- When value arrives from memory, check if guess was correct:
 - Correct: Save speculative work ⇒ performance gain
 - Incorrect: Discard speculative work ⇒ no harm (?)

Problem: Side Effects

Architectural Guarantee

Register values eventually match the result of in-order execution

Speculative Execution


CPU regularly performs incorrect calculations, then deletes mistakes

Is making + discarding mistakes the same as in-order execution?

The processor executed instructions that were not supposed to run !!


The problem: instructions can have observable side-effects

Spectre and Meltdown



Speculative execution bugs in Intel x86, ARM, IBM processors + cache-based side-channels (F+R)

- Allows reading kernel (or hypervisor, other VM) memory
- Similar attacks on SGX, etc.



Intel didn't warn US government about CPU security flaws until they were public

Meltdown and Spectre were kept secret

Researchers find malware samples that exploit Meltdown and Spectre

As of Feb. 1, antivirus testing firm AV-TEST had found 139 malware samples that exploit Meltdown and Spectre. Most are not very functional, but that could change.

Conditional Branch (Var 1) Attack

```
if (x < array1_size)
    y = array2[array1[x]*4096];
```

Suppose `unsigned int x` comes from untrusted caller

Execution without speculation is safe:

```
array2[array1[x]*4096] not eval unless x < array1_size
```

What about with speculative execution?

Conditional Branch (Var 1) Attack

```
if (x < array1_size)
    y = array2[array1[x]*4096];
```

Before attack:

- Train branch predictor to expect if() is true (e.g. call with `x < array1_size`)
- Evict `array1_size` and `array2[]` from cache

Memory & Cache Status

`array1_size = 00000008`

Memory at `array1` base:

8 bytes of data (value doesn't matter)

Memory at `array1` base+1000:

`09 F1 98 CC 90...` (something secret)

`array2[0*4096]`
`array2[1*4096]`
`array2[2*4096]`
`array2[3*4096]`
`array2[4*4096]`
`array2[5*4096]`
`array2[6*4096]`
`array2[7*4096]`
`array2[8*4096]`
`array2[9*4096]`
`array2[10*4096]`
`array2[11*4096]`

...

Contents don't matter
only care about cache
status

Uncached

Cached

Conditional Branch (Var 1) Attack

```
if (x < array1_size)
    y = array2[array1[x]*4096];
```

Attacker calls victim with $x=1000$

Speculative exec while waiting for `array1_size`:

- Predict that `if()` is true
- Read address (`array1 base + x`)
(using out-of-bounds $x=1000$)
- Read returns secret byte = **09**
(in cache \Rightarrow fast)

Memory & Cache Status

`array1_size = 00000008` ←

Memory at `array1` base:

8 bytes of data (value doesn't matter)

Memory at `array1` base+1000:

09 F1 98 CC 90... (something secret)

`array2[0*4096]`
`array2[1*4096]`
`array2[2*4096]`
`array2[3*4096]`
`array2[4*4096]`
`array2[5*4096]`
`array2[6*4096]`
`array2[7*4096]`
`array2[8*4096]`
`array2[9*4096]`
`array2[10*4096]`
`array2[11*4096]`
...

Contents don't matter
only care about cache
status

Uncached

Cached

Conditional Branch (Var 1) Attack

```
if (x < array1_size)
    y = array2[array1[x]*4096];
```

Attacker calls victim with $x=1000$

- Request mem at
(array2 base + 09*4096)
- Brings array2[09*4096] into the cache
- Realize if() is false,
discard speculative work

Finish operation & return to caller

Memory & Cache Status

array1_size = 00000008

Memory at array1 base:

8 bytes of data (value doesn't matter)

Memory at array1 base+1000:

09 F1 98 CC 90... (something secret)

array2[0*4096]
array2[1*4096]
array2[2*4096]
array2[3*4096]
array2[4*4096]
array2[5*4096]
array2[6*4096]
array2[7*4096]
array2[8*4096]
array2[9*4096]
array2[10*4096]
array2[11*4096]
...

Contents don't matter
only care about cache
status

Uncached

Cached

Conditional Branch (Var 1) Attack

```
if (x < array1_size)
    y = array2[array1[x]*4096];
```

Attacker calls victim with $x=1000$

- Measures read time for `array2[i*4096]`
- Read for $i=09$ is fast (cached!), reveals secret byte !!
- Repeat with many x (10KB/s)

Memory & Cache Status

`array1_size = 00000008`

Memory at `array1` base:

8 bytes of data (value doesn't matter)

Memory at `array1` base+1000:

09 F1 98 CC 90... (something secret)

```
array2[ 0*4096]
array2[ 1*4096]
array2[ 2*4096]
array2[ 3*4096]
array2[ 4*4096]
array2[ 5*4096]
array2[ 6*4096]
array2[ 7*4096]
array2[ 8*4096]
array2[ 9*4096]
array2[10*4096]
array2[11*4096]
...
```

Contents don't matter
only care about cache
status

Uncached

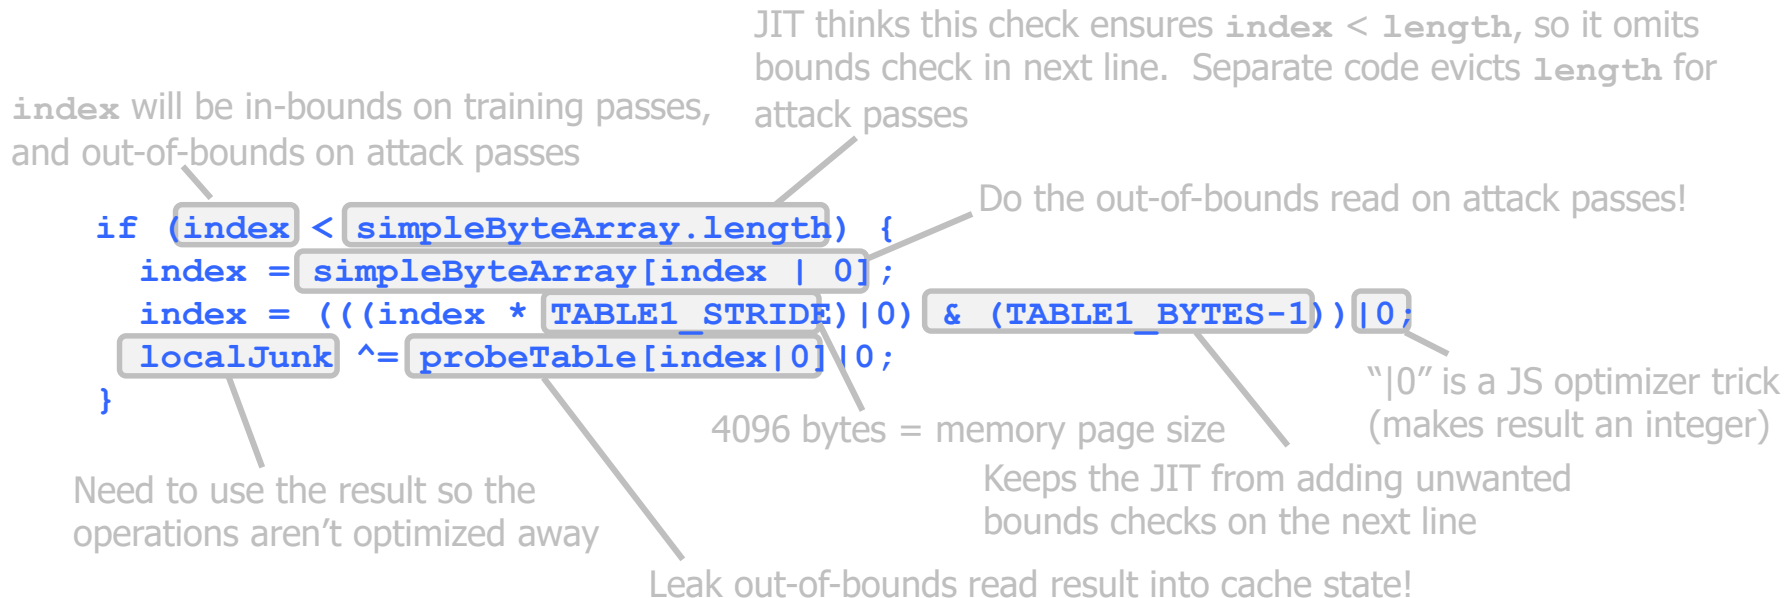
Cached

Violating JavaScript Sandbox

Browsers run JavaScript from untrusted websites

- JIT compiler inserts safety checks, including bounds checks on array accesses

Speculative execution runs through safety checks...



Can evict `length`/`probeTable` from JavaScript (easy)

... then use timing to detect newly-cached location in `probeTable`

Indirect Branches (Var 2)

Indirect branches can go anywhere, e.g., `jmp [rax]`

- If destination is delayed, CPU guesses and proceeds speculatively
- Find an indirect `jmp` with attacker-controlled register(s), then cause mispredict to a useful 'gadget'

```
y = array2[array1[x]*4096];
```

Attack steps:

- **Mistrain** branch prediction so speculative execution will go to gadget
- **Evict** address `[rax]` from cache to cause speculative execution
- **Execute** victim so it runs gadget speculatively
- **Detect** change in cache state to determine memory data

Mitigating Spectre

How to prevent Spectre without a huge cost in performance?

Idea 1: fully restore cache state when speculation fails

Insecure! Speculative execution can have observable side effects beyond the cache state

```
if (x < array1_size) {  
    y = array1[x];  
    do_something_observable(y);  
}
```

← occupy a bus (detectable from another core, or cause EM radiation)

Stopping Speculation

Idea: insert **LFENCE** on all vulnerable code paths

Efficient, no impact on benchmark software

```
if (x < array1_size)
    LFENCE           // processor instruction
    y = array2[ array1[x]*4096 ];
```

Insert  **LFENCES** manually?

Often millions of control flow paths

Too confusing - speculation runs 188++ instructions, crosses modules

Too risky – miss one and attacker can read entire process memory

Put  **LFENCES** everywhere?

Abysmal performance - **LFENCE** is very slow

Not in binary libraries, compiler-created code patterns

Insert  by smart compiler?

Protect only known-bad bad patterns = unsafe

- Microsoft Visual C/C++ /Qspectre unsafe for 13 of 15 tests

⇒ Protect all potentially-exploitable patterns

Transfer of blame (CPU -> SW): “you should have put an **LFENCE** there”

Remove All Branches?

DOOM with no branches:
one frame every ~ 7 hours

A branchless DOOM

This directory provides a branchless, mov-only version of the classic DOOM video game.



DOOM, running with only mov instructions.

This is thought to be entirely secure against the Meltdown and Spectre CPU vulnerabilities, which require speculative execution on branch instructions.

Oops! Idea 4: speculative store

More Attacks

Meltdown

Foreshadow

Rogue inflight data load (**RIDL**) and **Fallout**


ZombieLoad

Store-to-leak forwarding

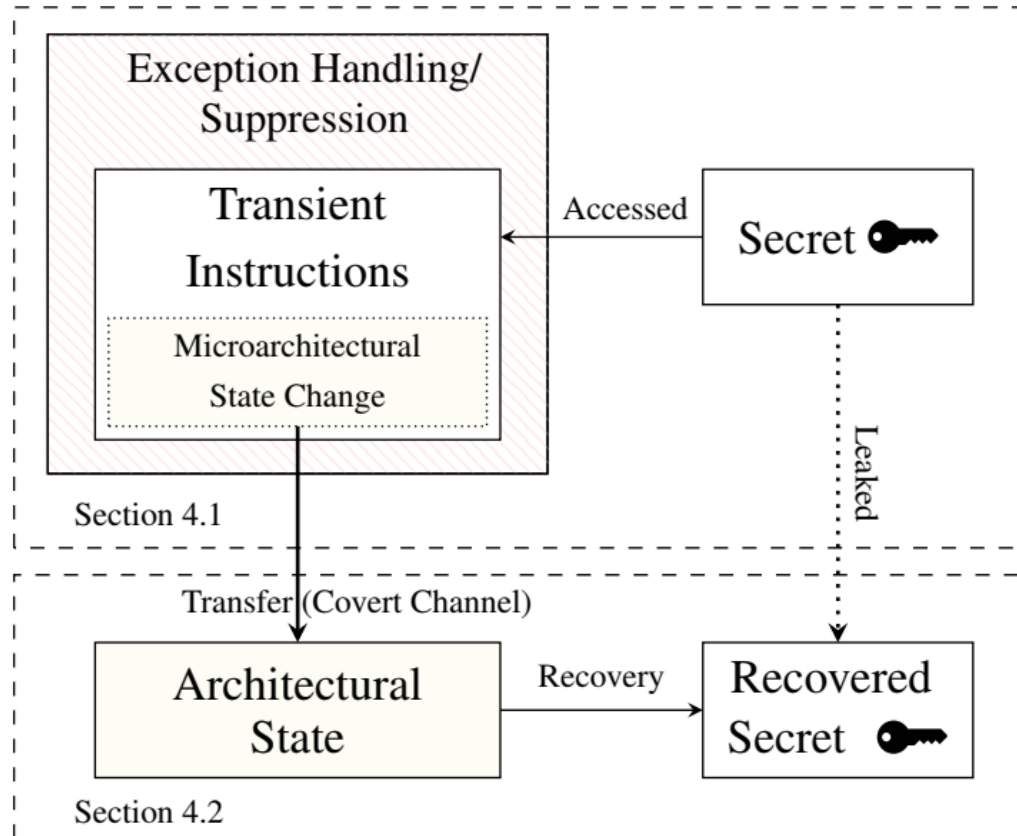
Enable reading unauthorized memory (client, cloud, SGX), mitigating incurs significant performance costs

Meltdown: Intuition

```
1 raise_exception();  
2 // the line below is never reached  
3 access(probe_array[data * 4096]);
```



Meltdown: Design



Meltdown: Core Spy Code

Retry reading privileged memory	→	1 ; rcx = kernel address
		2 ; rbx = probe array
		3 retry:
Access privileged memory	→	4 mov al, byte [rcx]
		5 shl rax, 0xc
Multiply by page size	→	6 jz retry
		7 mov rbx, qword [rbx + rax]

Read from an attacker
(unprivileged) array at:
 $(\text{secret value}) * 2^{12}$

Attacker times accessing $[\text{rbx} + \text{rax}]$ for different values of rax
When finds one that loads fast, learns sensitive byte